

---

# Machine Learning Laboratory Manual

Completed Edition with Full Implementations and Results

---

Practical Guide to Classical, Ensemble, Neural, Probabilistic,  
Density-Based, Semi-Supervised and Generative Machine Learning

Prepared by

**Dr. Ohidujjaman Tuhin**

Associate Professor, Dept. of CSE, UIU

Department of Computer Science & Engineering  
Machine Learning Lab

September 25, 2026

---

# Contents

|                                                          |               |
|----------------------------------------------------------|---------------|
| <b>Preface</b>                                           | <b>3</b>      |
| <b>Important Terminology Correction</b>                  | <b>4</b>      |
| <b>Scope of the Laboratory</b>                           | <b>5</b>      |
| <b>General Learning Outcomes</b>                         | <b>6</b>      |
| <b>Recommended Software Environment</b>                  | <b>7</b>      |
| <b>Laboratory Guidelines</b>                             | <b>8</b>      |
| <b>Minimum Report Format</b>                             | <b>9</b>      |
| <b>Recommended Evaluation Metrics</b>                    | <b>10</b>     |
| <br><b>I Clustering</b>                                  | <br><b>11</b> |
| Experiment 1: K-means Clustering                         | 12            |
| Experiment 2: Modified K-means                           | 16            |
| Experiment 3: Hierarchical Clustering                    | 19            |
| Experiment 4: Fuzzy C-means                              | 22            |
| <br><b>II Density-Based Learning</b>                     | <br><b>26</b> |
| Experiment 5: DBSCAN                                     | 27            |
| Experiment 6: HDBSCAN                                    | 31            |
| <br><b>III Semi-Supervised Learning</b>                  | <br><b>34</b> |
| Experiment 7: Self-training for Semi-supervised Learning | 35            |
| <br><b>IV Ensemble Learning</b>                          | <br><b>39</b> |
| Experiment 8: Random Forest Regression (RFR)             | 40            |
| Experiment 9: Random Forest Classification (RFC)         | 43            |
| Experiment 10: XGBoost                                   | 46            |
| Experiment 11: AdaBoost                                  | 49            |
| Experiment 12: CatBoost                                  | 52            |

|                                                                            |           |
|----------------------------------------------------------------------------|-----------|
| <b>V Neural Networks: Multilayer Perceptron</b>                            | <b>55</b> |
| Experiment 13: Multilayer Perceptron (MLP)                                 | 56        |
| <b>VI Sequence Modeling: Recurrent Neural Networks</b>                     | <b>61</b> |
| Experiment 14: Recurrent Neural Network (RNN)                              | 62        |
| <b>VII Unsupervised Neural Learning: Self-Organizing Maps</b>              | <b>68</b> |
| Experiment 15: Self-Organizing Map (SOM)                                   | 69        |
| <b>VIII Probabilistic Sequence Modeling: Hidden Markov Models</b>          | <b>73</b> |
| Experiment 16: Hidden Markov Model (HMM)                                   | 74        |
| <b>IX Margin-Based Learning: Support Vector Machines</b>                   | <b>78</b> |
| Experiment 17: Support Vector Machine (SVM)                                | 79        |
| <b>X Generative / Foundation Models: Large Language Models</b>             | <b>84</b> |
| Experiment 18: Large Language Model (LLM) Experiment                       | 85        |
| <b>XI Non-Parametric Regression: Generalized Regression Neural Network</b> | <b>90</b> |
| Experiment 19: Generalized Regression Neural Network (GRNN)                | 91        |
| <b>XII Integrated Mini Project and Appendices</b>                          | <b>92</b> |
| Integrated Mini Project: Comparative Machine Learning Study                | 93        |
| Appendix A: Quick Reference Formulas                                       | 94        |
| Appendix B: Reproducibility Checklist                                      | 95        |
| Appendix C: Suggested Assessment Rubric                                    | 96        |
| Appendix D: Student Experiment Record                                      | 97        |

## Preface

This manual is designed around the algorithm list supplied for the Machine Learning assignment. It turns that list into a structured laboratory sequence emphasizing implementation, experimentation, model evaluation and interpretation rather than only code execution. Students should run the programs, record outputs, vary important hyperparameters, and explain why the observed results change.

This **completed edition** goes one step further: every experiment is implemented and executed on real, public datasets, with the observed results, discussion and the complete runnable code included, so each laboratory can be reproduced standalone.

## Important Terminology Correction

The supplied list contains “Random Forest Regression (RFR)” and “Random Forest Regression (RFC)”. In standard machine-learning terminology, RFC normally means Random Forest Classifier. This manual therefore treats **RFR** as Random Forest Regressor and **RFC** as Random Forest Classifier.

## Scope of the Laboratory

| Module                         | Algorithms / Topics                                                                  |
|--------------------------------|--------------------------------------------------------------------------------------|
| Clustering                     | K-means, Modified K-means, Hierarchical Clustering, Fuzzy C-means                    |
| Density-based learning         | DBSCAN, HDBSCAN                                                                      |
| Semi-supervised learning       | Self-training                                                                        |
| Ensemble learning              | Random Forest Regression, Random Forest Classification, XGBoost, AdaBoost, Cat-Boost |
| Neural networks                | MLP, RNN (LSTM)                                                                      |
| Unsupervised neural learning   | SOM                                                                                  |
| Probabilistic sequence model   | HMM                                                                                  |
| Margin-based learning          | SVM (SVC and SVR)                                                                    |
| Generative / foundation models | LLM                                                                                  |
| Non-parametric regression      | GRNN                                                                                 |

## General Learning Outcomes

- Implement and execute major machine-learning algorithms using Python.
- Select an appropriate algorithm based on data type, task and assumptions.
- Preprocess data, handle categorical variables, scale numerical features and split datasets correctly.
- Evaluate clustering with internal measures and supervised models with appropriate metrics.
- Interpret model behavior through decision boundaries, feature importance, learning curves and error analysis.
- Compare algorithms using a fair experimental protocol rather than relying on a single score.
- Document reproducible experiments, hyperparameters, random seeds, observations and limitations.

## Recommended Software Environment

```
# Core
pip install numpy pandas matplotlib seaborn scikit-learn
# Specialized algorithms
pip install scikit-fuzzy hdbscan hmmlearn minisom
# Gradient boosting libraries
pip install xgboost catboost
# LLM / transformer experiment
pip install transformers torch
```

**Listing 1:** Core and specialized libraries used by the manual



## Laboratory Guidelines

| Before Lab                                                                                                              | During Lab                                                                                                                               | After Lab                                                                                                          |
|-------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------|
| Read theory and algorithm steps.<br>Know the dataset and target.<br>Predict expected behavior.<br>Prepare key formulas. | Run the baseline program.<br>Record hyperparameters and seed.<br>Perform at least one controlled variation.<br>Check train/test leakage. | Submit code and results.<br>Explain observations.<br>Answer viva questions.<br>State limitations and improvements. |

## Minimum Report Format

1. Experiment title and student information.
2. Problem statement and objective.
3. Dataset description and preprocessing.
4. Algorithm principle and pseudocode / flowchart.
5. Python implementation.
6. Hyperparameters and experimental setup.
7. Results: tables, plots and relevant metrics.
8. Discussion: what the results mean and why they occurred.
9. Conclusion and limitations.
10. Answers to viva questions.

## Recommended Evaluation Metrics

| Task                | Suggested Metrics                                                             |
|---------------------|-------------------------------------------------------------------------------|
| Classification      | Accuracy, Precision, Recall, F1-score, ROC-AUC, confusion matrix              |
| Regression          | MAE, MSE, RMSE, $R^2$                                                         |
| Clustering          | Silhouette score, Calinski–Harabasz, Davies–Bouldin; visual inspection        |
| Sequence modeling   | Negative log-likelihood / log-likelihood, sequence accuracy where appropriate |
| Language generation | Task-specific automatic metrics plus qualitative / error analysis             |

# **Part I**

## **Clustering**

## Experiment 1: K-means Clustering

**Dataset:** Mall Customers (`data/mall_customers.csv`); 200 customers, 2 features (Annual Income, Spending Score)

**Key parameters:**  $K = 5$ , K-Means++ initialization, `n_init = 20`, `StandardScaler`

**Headline result:** Silhouette 0.5547; Davies–Bouldin 0.5722; Calinski–Harabasz 248.65

### Objective

Partition unlabeled observations into  $K$  clusters by iteratively assigning points to the nearest centroid and updating centroids; choose  $K$  with the elbow and silhouette diagnostics.

### Suggested Dataset

Mall Customers (`data/mall_customers.csv`), using Annual Income and Spending Score as a two-dimensional, visualizable feature space. Features are standardized (`StandardScaler`).

### Required Libraries

numpy, pandas, matplotlib, scikit-learn.

### Brief Theory

K-means minimizes the within-cluster sum of squared distances

$$J = \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2,$$

by alternating two steps: assign every point to the nearest centroid, then set each centroid to the mean of its assigned points. Each iteration never increases  $J$ , so the algorithm converges — but only to a *local* optimum, which is why initialization matters. **K-Means++** spreads the initial centroids by sampling them with probability proportional to the squared distance to the nearest chosen centroid, giving lower and more consistent inertia. Because Euclidean distance is scale-sensitive, numerical features should be standardized.  $K$  itself must be chosen externally: the *elbow* (inertia) and the *silhouette* (cohesion vs. separation) are the standard diagnostics.

### Procedure

1. Load the dataset and retain the two numerical features.
2. Standardize the features.
3. Try  $K = 2 \dots 10$ ; record inertia (WCSS) and silhouette for each  $K$ .
4. Choose  $K$  at the elbow / silhouette peak and fit the final model with K-Means++ initialization, 20 restarts and a fixed random seed (see the code listing).
5. Visualize the clusters and centroids.
6. Report silhouette, Davies–Bouldin and Calinski–Harabasz.

### Reference Implementation

```
for k in range(2, 11):
    kmeans = KMeans(n_clusters=k, init='k-means++', n_init=10, random_state=42)
    kmeans.fit(X_mall_scaled)
    wcss.append(kmeans.inertia_)
    silhouette_scores.append(silhouette_score(X_mall_scaled, kmeans.labels_))


kmeans_model = KMeans(n_clusters=5, init='k-means++', n_init=20, random_state=42)
kmeans_labels = kmeans_model.fit_predict(X_mall_scaled)
```

---

**Listing 2:** Selecting  $K$ , then fitting the final K-means model

## Evaluation / Expected Output

Both the elbow curve and the silhouette curve indicate  $K = 5$  (Figure 1).



**Figure 1:** Elbow (inertia) and silhouette curves for  $K = 2 \dots 10$ ; the elbow and the silhouette peak both occur at  $K = 5$ .

**Table 4:** K-means evaluation on the standardized Mall Customers features.

| Metric                                 | Value         |
|----------------------------------------|---------------|
| Selected $K$ (elbow + silhouette peak) | 5             |
| Silhouette score                       | <b>0.5547</b> |
| Davies–Bouldin index                   | 0.5722        |
| Calinski–Harabasz score                | 248.65        |

Silhouette  $\approx 0.55$  indicates compact, reasonably separated clusters on real data; Davies–Bouldin  $< 1$  confirms good separation.

## Viva / Discussion Questions

**Why does K-means depend on feature scaling?** It minimizes Euclidean distance; unscaled features with larger ranges dominate the distance and therefore the assignments.

**What is inertia?** The within-cluster sum of squared distances to the centroids (WCSS); lower is tighter, but it always decreases with  $K$ , so it is not a standalone quality score.

**Why can a high silhouette score still be misleading?** It favors compact convex clusters and can be high for a trivial partition; it misjudges elongated or density-varying structures.

**What happens when clusters are non-spherical?** K-means splits them with straight Voronoi boundaries and merges elongated clusters incorrectly (see Experiment 5).

**Why is K-Means++ preferred over random initialization?** It spreads the initial centroids out, reducing convergence to poor local optima and giving more stable inertia.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Scikit-Learn clustering models ----
```

```

from sklearn.cluster import KMeans, AgglomerativeClustering
# ---- Feature scaling + unsupervised evaluation metrics ----
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import silhouette_score, davies_bouldin_score, calinski_harabasz_score
# ---- SciPy hierarchical clustering utilities (linkage + dendrogram) ----
from scipy.cluster.hierarchy import dendrogram, linkage

# Styling setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110

# Fix the random seed so every run reproduces the same results
np.random.seed(42)
print("Libraries successfully imported!")

```

**Listing 3:** Listing 1.1 Imports and plotting style

```

# ---- Load the Mall Customers dataset ----
df_mall = pd.read_csv('../data/mall_customers.csv')
print("Dataset Shape:", df_mall.shape)
print(df_mall.head(3))

# ---- Feature selection: two features so the clusters can be shown in 2D ----
X_mall = df_mall[['Annual Income (k$)', 'Spending Score (1-100)']].values

# ---- Standardize (mean = 0, std = 1): distance-based algorithms are scale-sensitive ----
scaler = StandardScaler()
X_mall_scaled = scaler.fit_transform(X_mall)

```

**Listing 4:** Listing 1.2 Load and preprocess the Mall Customers data

```

# ---- Selecting the optimal number of clusters K ----
# Try K = 2..10 and record two diagnostics:
# * Inertia (WCSS): total squared distance from points to their centroid
# * Silhouette score: cluster cohesion vs. separation
k_range = range(2, 11)
wcss = []
silhouette_scores = []

for k in k_range:
    # n_init=10 runs K-Means from 10 different seeds and keeps the best run
    kmeans = KMeans(n_clusters=k, init='k-means++', n_init=10, random_state=42)
    kmeans.fit(X_mall_scaled)
    wcss.append(kmeans.inertia_) # WCSS of the fitted model
    silhouette_scores.append(silhouette_score(X_mall_scaled, kmeans.labels_)) # partition quality

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))

# ---- Left: elbow curve (look for the "knee" where inertia stops falling quickly) ----
ax1.plot(k_range, wcss, 'bo-', linewidth=2, markersize=8)
ax1.set_title("Elbow Method for Optimal K", fontsize=14, fontweight='bold')
ax1.set_xlabel("Number of Clusters (K)", fontsize=12)
ax1.set_ylabel("Inertia (WCSS)", fontsize=12)
ax1.axvline(x=5, color='r', linestyle='--', label='Elbow at K=5')
ax1.legend()

# ---- Right: silhouette score for each K (higher is better) ----
ax2.plot(k_range, silhouette_scores, 'go-', linewidth=2, markersize=8)
ax2.set_title("Silhouette Scores across K", fontsize=14, fontweight='bold')
ax2.set_xlabel("Number of Clusters (K)", fontsize=12)
ax2.set_ylabel("Silhouette Score", fontsize=12)
ax2.axvline(x=5, color='r', linestyle='--', label='Optimal K=5')
ax2.legend()

plt.tight_layout()
plt.show()

```

**Listing 5:** Listing 1.3 Select K with the elbow and silhouette curves

```

# ---- Final K-Means model with the selected K = 5 ----
optimal_k = 5

```

```
kmeans_model = KMeans(n_clusters=optimal_k, init='k-means++', n_init=20, random_state=42)
kmeans_labels = kmeans_model.fit_predict(X_mall_scaled) # hard cluster assignment per sample

print(f"K-Means Silhouette Score: {silhouette_score(X_mall_scaled, kmeans_labels):.4f}")
print(f"K-Means Davies-Bouldin Index: {davies_bouldin_score(X_mall_scaled, kmeans_labels):.4f}")
```

**Listing 6:** Listing 1.4 Final K-means model (K = 5)



## Experiment 2: Modified K-means

**Dataset:** Mall Customers ( $200 \times 2$ , standardized)

**Key parameters:** K-Means++ initialization; outlier threshold = 97th percentile of centroid distances

**Headline result:** Distance threshold 1.072; 6 of 200 observations flagged as candidate outliers

### Objective

Implement a transparent modification of K-means that uses K-Means++ initialization and an explicit outlier-distance check before final cluster interpretation.

### Suggested Dataset

Mall Customers ( $200 \times 2$ , standardized). The modification is applied on top of the same features.

### Required Libraries

numpy, matplotlib, scikit-learn.

### Brief Theory

There is no single universally accepted algorithm named “modified K-means”; variants differ by application. In this lab the modification is intentionally defined so that it can be reproduced exactly:

1. initialize the centroids with **K-Means++** (spread-out seeds);
2. fit standard K-means and obtain the hard labels;
3. compute each sample’s distance to its *assigned* centroid,  $d_i = \|x_i - \mu_{c_i}\|$ ;
4. set a threshold from a documented percentile of the training distances (97th percentile here);
5. flag samples with  $d_i$  beyond the threshold as **candidate outliers**.

This is an experimental variant that adds outlier awareness to the standard algorithm; it is not a standard replacement for robust clustering methods.

### Procedure

1. Fit K-means with K-Means++ initialization,  $K = 5$ .
2. Compute  $d_i = \|x_i - \mu_{c_i}\|$  for every sample.
3. Set the threshold to the 97th percentile of the training distances.
4. Flag samples with  $d_i > \text{threshold}$  as candidate outliers.
5. Plot the clusters with the flagged points highlighted.
6. Compare the modified interpretation with the standard labels.

### Reference Implementation

```
km_mod = KMeans(n_clusters=5, init='k-means++', n_init=20, random_state=42)
mod_labels = km_mod.fit_predict(X_mall_scaled)

dist_to_centroid = np.linalg.norm(X_mall_scaled - km_mod.cluster_centers_[mod_labels], axis=1)
threshold = np.percentile(dist_to_centroid, 97)
outlier_flag = dist_to_centroid > threshold
```

**Listing 7:** K-Means++ followed by the outlier-distance check

## Evaluation / Expected Output

**Table 5:** Modified K-means: outlier-distance check on the standardized features.

| Item                                         | Value                         |
|----------------------------------------------|-------------------------------|
| Distance threshold (97th percentile)         | 1.072 (standardized units)    |
| Flagged candidate outliers                   | <b>6 of 200 (3%)</b>          |
| Cluster labels                               | identical to standard K-means |
| Silhouette / Davies–Bouldin (same partition) | 0.5547 / 0.5722               |

The flagged customers are the most weakly attached points of their clusters — typically boundary customers with unusual income/spending combinations. The cluster structure itself is unchanged, which is expected: the modification adds an outlier interpretation layer.

## Viva / Discussion Questions

**Why is it important to define a modification precisely?** Without a formal definition “modified K-means” is ambiguous; fixing the initialization, the distance measure and the percentile rule makes the experiment reproducible and testable.

**How does K-Means++ differ from random initialization?** K-Means++ selects initial centroids with probability proportional to the squared distance to the nearest chosen centroid, spreading the seeds out and avoiding poor local optima.

**Why might distance-to-centroid fail for elongated clusters?** Centroid distance is isotropic; a point can be close to the centroid along the long axis yet far along the short axis, so elongated clusters can flag legitimate points and hide real outliers.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Scikit-Learn clustering models ----
from sklearn.cluster import KMeans, AgglomerativeClustering
# ---- Feature scaling + unsupervised evaluation metrics ----
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import silhouette_score, davies_bouldin_score, calinski_harabasz_score
# ---- SciPy hierarchical clustering utilities (linkage + dendrogram) ----
from scipy.cluster.hierarchy import dendrogram, linkage

# Styling setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110

# Fix the random seed so every run reproduces the same results
np.random.seed(42)
print("Libraries successfully imported!")
```

**Listing 8:** Core numerical and plotting libraries

```
# ---- Load the Mall Customers dataset ----
df_mall = pd.read_csv('../data/mall_customers.csv')
print("Dataset Shape:", df_mall.shape)
print(df_mall.head(3))

# ---- Feature selection: two features so the clusters can be shown in 2D ----
X_mall = df_mall[['Annual Income (k$)', 'Spending Score (1-100)']].values

# ---- Standardize (mean = 0, std = 1): distance-based algorithms are scale-sensitive ----
```

```
scaler = StandardScaler()
X_mall_scaled = scaler.fit_transform(X_mall)
```

**Listing 9:** Load the Mall Customers dataset

```
# ---- Modified K-Means (manual variant): K-Means++ + outlier-distance check ----
km_mod = KMeans(n_clusters=optimal_k, init='k-means++', n_init=20, random_state=42)
mod_labels = km_mod.fit_predict(X_mall_scaled)

# Distance of every sample to its assigned centroid
dist_to_centroid = np.linalg.norm(X_mall_scaled - km_mod.cluster_centers_[mod_labels], axis=1)
threshold = np.percentile(dist_to_centroid, 97)      # documented percentile
outlier_flag = dist_to_centroid > threshold

print(f"Distance threshold (97th percentile): {threshold:.3f}")
print(f"Flagged candidate outliers: {outlier_flag.sum()} of {len(X_mall)}")

# ---- Plot clusters with flagged candidate outliers highlighted ----
plt.figure(figsize=(9, 6))
for cluster_id in range(optimal_k):
    pts = X_mall[mod_labels == cluster_id]
    plt.scatter(pts[:, 0], pts[:, 1], s=45, edgecolors='k', alpha=0.8)
plt.scatter(X_mall[outlier_flag, 0], X_mall[outlier_flag, 1], s=180, facecolors='none',
            edgecolors='red', linewidths=2.5, label=f'Flagged outliers ({outlier_flag.sum()})')
plt.title("Modified K-Means: K-Means++ with Outlier-Distance Flagging", fontsize=13, fontweight='bold')
plt.xlabel("Annual Income (k$)", fontsize=12)
plt.ylabel("Spending Score (1-100)", fontsize=12)
plt.legend()
plt.tight_layout()
plt.show()
```

**Listing 10:** Modified K-Means (manual variant): K-Means++ + outlier-distance check

## Experiment 3: Hierarchical Clustering

**Dataset:** Mall Customers ( $200 \times 2$ , standardized)

**Key parameters:** Ward linkage; cut at  $K = 5$ ; truncated dendrogram (last 25 merges)

**Headline result:** Silhouette 0.5538; Davies–Bouldin 0.5779; Calinski–Harabasz 244.41

### Objective

Perform agglomerative hierarchical clustering and interpret the resulting dendrogram and cluster labels.

### Suggested Dataset

Mall Customers ( $200 \times 2$ , standardized).

### Required Libraries

scipy, scikit-learn, matplotlib.

### Brief Theory

Agglomerative clustering starts with one cluster per point and repeatedly merges the closest pair of clusters. With **Ward linkage** the merge cost is the increase in total within-cluster variance,

$$\Delta W(A, B) = \frac{n_A n_B}{n_A + n_B} \|\mu_A - \mu_B\|^2,$$

so the algorithm always merges the pair whose combination loses the least compactness. The dendrogram records every merge and its height (distance); cutting it at a chosen height yields a flat partition with the desired number of clusters.

### Procedure

1. Standardize the features.
2. Compute the Ward linkage matrix.
3. Plot the dendrogram (truncated for readability) and choose a cut level.
4. Fit `AgglomerativeClustering` with  $K = 5$  and Ward linkage.
5. Evaluate with silhouette, Davies–Bouldin and Calinski–Harabasz.
6. Discuss the choice of linkage.

### Reference Implementation

```
linkage_matrix = linkage(X_mall_scaled, method='ward')
dendrogram(linkage_matrix, truncate_mode='lastp', p=25, show_contracted=True)


agg_model = AgglomerativeClustering(n_clusters=5, metric='euclidean', linkage='ward')
agg_labels = agg_model.fit_predict(X_mall_scaled)
```

**Listing 11:** Ward linkage, dendrogram and the final agglomerative model

### Evaluation / Expected Output

The dendrogram (Figure 2) shows a clear large merge gap that supports cutting at five clusters.

The partition is nearly identical in quality to K-means (silhouette difference below 0.001), but requires no initialization and exposes the full merge hierarchy.



**Figure 2:** Hierarchical clustering dendrogram (Ward linkage, truncated to the last 25 merges); the dashed line marks the cut that yields  $K = 5$ .

**Table 6:** Hierarchical (Ward) clustering evaluation on the standardized features.

| Metric                  | Value  |
|-------------------------|--------|
| Silhouette score        | 0.5538 |
| Davies–Bouldin index    | 0.5779 |
| Calinski–Harabasz score | 244.41 |

## Viva / Discussion Questions

**What does a dendrogram show?** The complete merge hierarchy: leaf order, merge heights (distances) and the nested cluster structure; cutting at a height produces a flat partition.

**What is linkage?** The rule that defines the distance between two clusters: single (closest pair), complete (farthest pair), average (mean pair) or Ward (variance increase).

**When is Ward linkage inappropriate?** With strongly non-spherical/elongated clusters, very unequal cluster sizes or heavy outliers — it assumes compact clusters of similar variance.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Scikit-Learn clustering models ----
from sklearn.cluster import KMeans, AgglomerativeClustering
# ---- Feature scaling + unsupervised evaluation metrics ----
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import silhouette_score, davies_bouldin_score, calinski_harabasz_score
# ---- SciPy hierarchical clustering utilities (linkage + dendrogram) ----
from scipy.cluster.hierarchy import dendrogram, linkage

# Styling setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110

# Fix the random seed so every run reproduces the same results
np.random.seed(42)
print("Libraries successfully imported!")
```

**Listing 12:** Core numerical and plotting libraries

```
# ---- Load the Mall Customers dataset ----
df_mall = pd.read_csv('../data/mall_customers.csv')
print("Dataset Shape:", df_mall.shape)
print(df_mall.head(3))

# ---- Feature selection: two features so the clusters can be shown in 2D ----
X_mall = df_mall[['Annual Income (k$)', 'Spending Score (1-100)']].values

# ---- Standardize (mean = 0, std = 1): distance-based algorithms are scale-sensitive ----
scaler = StandardScaler()
X_mall_scaled = scaler.fit_transform(X_mall)
```

**Listing 13:** Load the Mall Customers dataset

```
# ---- Build the linkage matrix and draw the dendrogram ----
# 'ward' linkage merges the pair of clusters that increases within-cluster variance the least.
linkage_matrix = linkage(X_mall_scaled, method='ward')

# The dendrogram shows every merge; we truncate it to the last 25 merges for readability.
plt.figure(figsize=(12, 5))
dendrogram(
    linkage_matrix,
    truncate_mode='lastp',    # show only the last p merged clusters
    p=25,
    leaf_rotation=45,
    leaf_font_size=10,
    show_contracted=True     # condensed leaves are shown as counts
)
plt.title("Hierarchical Clustering Dendrogram (Ward Linkage)", fontsize=14, fontweight='bold')
plt.xlabel("Cluster Sample Index / Merged Size", fontsize=12)
plt.ylabel("Euclidean Distance (Ward Threshold)", fontsize=12)
plt.axhline(y=10.0, color='r', linestyle='--', label='Cut-off Threshold (K=5)') # horizontal cut => cluster count
plt.legend()
plt.tight_layout()
plt.show()
```

**Listing 14:** Build the linkage matrix and draw the dendrogram

```
# ---- Fit the final Agglomerative model (Ward linkage, K = 5) ----
agg_model = AgglomerativeClustering(n_clusters=optimal_k, metric='euclidean', linkage='ward')
agg_labels = agg_model.fit_predict(X_mall_scaled)

print(f"Hierarchical Silhouette Score: {silhouette_score(X_mall_scaled, agg_labels):.4f}")
print(f"Hierarchical Davies-Bouldin Index: {davies_bouldin_score(X_mall_scaled, agg_labels):.4f}")
```

**Listing 15:** Fit the final Agglomerative model (Ward linkage, K = 5)

## Experiment 4: Fuzzy C-means

**Dataset:** Mall Customers ( $200 \times 2$ , standardized);  $K = 5$   
**Key parameters:** Fuzzifier  $m = 2.0$ ; tolerance  $10^{-5}$ ; maximum 150 iterations  
**Headline result:** Silhouette 0.5547; Fuzzy Partition Coefficient (FPC) 0.6711

### Objective

Cluster data using fuzzy memberships so that each observation can belong to multiple clusters with different degrees of membership.

### Suggested Dataset

Mall Customers ( $200 \times 2$ , standardized),  $K = 5$ , fuzzifier  $m = 2.0$ .

### Required Libraries

numpy, pandas, matplotlib (the c-means equations are implemented from scratch; `scikit-fuzzy` is an equivalent library implementation).

### Brief Theory

Unlike hard K-means, Fuzzy C-means assigns every sample  $x_i$  a membership degree  $u_{ik} \in [0, 1]$  in every cluster, with  $\sum_{k=1}^K u_{ik} = 1$ . It minimizes

$$J_m = \sum_{i=1}^N \sum_{k=1}^K u_{ik}^m \|x_i - v_k\|^2, \quad m > 1,$$

by alternating two updates until convergence:

$$v_k = \frac{\sum_i u_{ik}^m x_i}{\sum_i u_{ik}^m}, \quad u_{ik} = \frac{1}{\sum_{j=1}^K \left( \frac{d_{ik}}{d_{ij}} \right)^{2/(m-1)}}.$$

The fuzzifier  $m$  controls how soft the memberships are ( $m \rightarrow 1$  approaches K-means). The **Fuzzy Partition Coefficient**  $\text{FPC} = \frac{1}{N} \sum_i \sum_k u_{ik}^2$  measures the fuzziness of the result (1 = completely hard,  $1/K$  = maximally fuzzy).

### Procedure

1. Standardize the data.
2. Initialize the membership matrix randomly, with each row summing to 1.
3. Alternate the centroid and membership updates until  $\max |\Delta U| < 10^{-5}$ .
4. Derive hard labels with  $\arg \max_k u_{ik}$ .
5. Report the clustering metrics and the FPC; visualize the hard partition.
6. Identify ambiguous samples with near-uniform membership vectors.

### Reference Implementation

```
Um = U ** self.m
centers = np.dot(Um.T, X) / Um.sum(axis=0)[:, np.newaxis]
dist = np.linalg.norm(X[:, np.newaxis, :] - centers[np.newaxis, :, :], axis=2) + 1e-10
U = (1.0 / dist) ** (2.0 / (self.m - 1.0))
U = U / U.sum(axis=1, keepdims=True)
labels = np.argmax(U, axis=1)
```


**Listing 16:** Core alternating updates of Fuzzy C-means

## Evaluation / Expected Output

**Table 7:** Fuzzy C-means evaluation ( $K = 5$ ,  $m = 2.0$ ).

| Metric                            | Value         |
|-----------------------------------|---------------|
| Silhouette score (hard labels)    | 0.5547        |
| Davies–Bouldin index              | 0.5722        |
| Fuzzy Partition Coefficient (FPC) | <b>0.6711</b> |
| Iterations to convergence         | fewer than 50 |

FPC = 0.67 (between  $1/K = 0.2$  and 1.0) shows that the partition is mostly crisp but retains soft boundary memberships. Figure 3 compares the four clustering partitions obtained in this part of the manual.



**Figure 3:** Clustering partitions on the Mall Customers data: K-means, modified K-means (outlier flagging), hierarchical Ward and Fuzzy C-means.

## Viva / Discussion Questions

**How is Fuzzy C-means different from K-means?** It assigns membership degrees instead of hard labels, so each point belongs to all clusters with weights that sum to 1; K-means is the limiting hard case.

**What does  $m$  control?** The fuzziness exponent: larger  $m$  spreads membership across clusters (softer, lower FPC), while  $m \rightarrow 1$  gives crisp K-means-like assignments.

**What does a membership vector represent?** The degree to which one observation belongs to each cluster; near-uniform values indicate an ambiguous point on a boundary.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import os
```



```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Scikit-Learn clustering models ----
from sklearn.cluster import KMeans, AgglomerativeClustering
# ---- Feature scaling + unsupervised evaluation metrics ----
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import silhouette_score, davies_bouldin_score, calinski_harabasz_score
# ---- SciPy hierarchical clustering utilities (linkage + dendrogram) ----
from scipy.cluster.hierarchy import dendrogram, linkage

# Styling setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110

# Fix the random seed so every run reproduces the same results
np.random.seed(42)
print("Libraries successfully imported!")

```

Listing 17: Core numerical and plotting libraries

```

# ---- Load the Mall Customers dataset ----
df_mall = pd.read_csv('../data/mall_customers.csv')
print("Dataset Shape:", df_mall.shape)
print(df_mall.head(3))

# ---- Feature selection: two features so the clusters can be shown in 2D ----
X_mall = df_mall[['Annual Income (k$)', 'Spending Score (1-100)']].values

# ---- Standardize (mean = 0, std = 1): distance-based algorithms are scale-sensitive ----
scaler = StandardScaler()
X_mall_scaled = scaler.fit_transform(X_mall)

```

Listing 18: Load the Mall Customers dataset

```

class FuzzyCMeans:
    """
    Fuzzy C-Means (FCM) clustering - implemented from scratch with NumPy.

    Each sample belongs to every cluster with a membership degree  $u_{ik}$  in  $[0, 1]$ ,
    and the memberships over the  $K$  clusters sum to 1 per sample.

    Parameters
    -----
    n_clusters    : number of fuzzy clusters  $K$ 
    m             : fuzzifier exponent ( $m > 1$ ; larger  $m \Rightarrow$  fuzzier memberships)
    max_iter      : maximum number of alternating-optimization iterations
    tol           : convergence tolerance on the maximum membership change
    random_state  : seed for reproducible random initialization
    """

    def __init__(self, n_clusters=5, m=2.0, max_iter=150, tol=1e-5, random_state=42):
        self.n_clusters = n_clusters
        self.m = m
        self.max_iter = max_iter
        self.tol = tol
        self.random_state = random_state

    def fit(self, X):
        rng = np.random.RandomState(self.random_state)
        n_samples = X.shape[0]

        # Step 1: initialize the membership matrix  $U$  randomly and normalize every row to sum to 1
        U = rng.rand(n_samples, self.n_clusters)
        U = U / U.sum(axis=1, keepdims=True)

        # Step 2: alternate between centroid and membership updates until convergence
        for iteration in range(self.max_iter):
            U_old = U.copy()

```

```

# 2a: centroid update - membership-weighted mean of the data
Um = U ** self.m                                     # u_ik^m
centers = np.dot(Um.T, X) / Um.sum(axis=0)[:, np.newaxis]

# 2b: compute the (n_samples x n_clusters) distance matrix to every centroid
dist = np.linalg.norm(X[:, np.newaxis, :] - centers[np.newaxis, :, :], axis=2)
dist = np.fmax(dist, 1e-10) # avoid division by zero

# 2c: membership update - closer points receive higher membership
#     u_ik = 1 / sum_j (d_ik / d_ij)^(2/(m-1))
inv_dist = 1.0 / dist
power = 2.0 / (self.m - 1.0)
inv_dist_p = inv_dist ** power
U = inv_dist_p / inv_dist_p.sum(axis=1, keepdims=True)

# 2d: stop when memberships stop changing (convergence)
if np.max(np.abs(U - U_old)) < self.tol:
    break

# Store results: final centroids, full membership matrix, and hard labels
self.cluster_centers_ = centers
self.u_ = U
self.labels_ = np.argmax(U, axis=1) # hard assignment = cluster with maximum membership
return self

# Fit FCM with the same K as the other algorithms (m = 2 is the standard fuzzifier)
fcm = FuzzyCMeans(n_clusters=optimal_k, m=2.0, random_state=42)
fcm.fit(X_small_scaled)
fcm_labels = fcm.labels_
fcm_centers = scaler.inverse_transform(fcm.cluster_centers_)

# ---- Evaluation of the hard partition induced by the fuzzy memberships ----
print(f"FCM Silhouette Score: {silhouette_score(X_small_scaled, fcm_labels):.4f}")
print(f"FCM Davies-Bouldin Index: {davies_bouldin_score(X_small_scaled, fcm_labels):.4f}")
fpc = np.mean(np.sum(fcm.u_ ** 2, axis=1)) # fuzzy partition coefficient (1/K fuzzy ... 1 hard)
print(f"FCM Fuzzy Partition Coefficient (FPC): {fpc:.4f}")

```

Listing 19: Fuzzy C-Means class (from scratch)

# **Part II**

## **Density-Based Learning**

## Experiment 5: DBSCAN

**Dataset:** USGS earthquakes 2023 (`data/earthquakes.csv`); 7,638 events with magnitude  $\geq 4.5$

**Key parameters:**  $\varepsilon = 0.03$  rad ( $\approx 191$  km), `min_samples = 10`, haversine metric on radian coordinates

**Headline result:** 55 clusters, 13.5% noise; largest clusters match the Philippines, Papua New Guinea, Tonga, Japan and Turkey

### Objective

Detect dense groups and noise using DBSCAN without specifying the number of clusters in advance, applied to real spatial data.

### Suggested Dataset

Real USGS earthquake catalogue 2023 (`data/earthquakes.csv`): 16,190 global events at magnitude  $\geq 4.0$ , filtered to the 7,638 events at magnitude  $\geq 4.5$ . Coordinates are converted to radians and clustered with the haversine (great-circle) metric.

### Required Libraries

scikit-learn, numpy, pandas, matplotlib.

### Brief Theory

DBSCAN defines clusters as connected regions of density. For a point  $p$ :

- **$\varepsilon$ -neighborhood:**  $N_\varepsilon(p) = \{q : \text{dist}(p, q) \leq \varepsilon\}$ ;
- **core point:**  $|N_\varepsilon(p)| \geq \text{min\_samples}$ ;
- **border point:** inside the  $\varepsilon$ -neighborhood of a core point but not core itself;
- **noise:** neither core nor border (label  $-1$ ).

Clusters grow by chaining density-reachable core points, so arbitrary shapes can be discovered and no  $K$  is required. A single global  $\varepsilon$  struggles when densities vary. Because latitude/longitude are spherical coordinates, the **haversine** (great-circle) metric is used with radian inputs;  $\varepsilon$  is then expressed in radians ( $0.03 \text{ rad} \times 6371 \text{ km} \approx 191 \text{ km}$ ).

### Procedure

1. Load the catalogue and filter to magnitude  $\geq 4.5$  (well-recorded events).
2. Convert latitude/longitude to radians for the haversine metric.
3. Choose  $\varepsilon$  from the  $k$ -distance graph ( $k = 10$ ); the knee is at 0.03 rad.
4. Fit DBSCAN and count clusters and noise events.
5. Show a parameter-sensitivity table for at least two settings.
6. Visualize the epicenters and validate the largest clusters against known seismic regions.

### Reference Implementation

```
X_rad = np.radians(quakes[['latitude', 'longitude']].values)
nbrs = NearestNeighbors(n_neighbors=10, metric='haversine').fit(X_rad)
k_distances = np.sort(nbrs.kneighbors(X_rad)[0][:, -1]) # knee ~0.03 rad

dbscan = DBSCAN(eps=0.03, min_samples=10, metric='haversine')
db_labels = dbscan.fit_predict(X_rad)
```

**Listing 20:**  $\varepsilon$  from the  $k$ -distance knee, then DBSCAN with the haversine metric


## Evaluation / Expected Output

**Table 8:** DBSCAN parameter sensitivity ( $\text{min\_samples} = 10$ ).

| $\varepsilon$ (rad)    | approx. km | clusters  | noise        |
|------------------------|------------|-----------|--------------|
| 0.02                   | 127        | 70        | 21.2%        |
| <b>0.03 (selected)</b> | <b>191</b> | <b>55</b> | <b>13.5%</b> |
| 0.05                   | 319        | 50        | 5.7%         |

**Table 9:** Largest DBSCAN clusters and their dominant region (real-world validation).

| Cluster | Events | Center lat/lon | Dominant region        |
|---------|--------|----------------|------------------------|
| 4       | 1,629  | 6.1, 126.1     | Philippines            |
| 0       | 1,068  | -12.5, 160.0   | Papua New Guinea       |
| 1       | 947    | -24.0, -154.5  | Tonga                  |
| 5       | 895    | 30.8, 145.0    | Japan region           |
| 2       | 253    | -57.8, -26.3   | South Sandwich Islands |
| 8       | 207    | 37.7, 37.2     | Turkey                 |



**Figure 4:** The  $k$ -distance graph ( $k = 10$ ); the knee at  $\varepsilon = 0.03$  rad marks the transition from dense seismic zones to sparse isolated events.

With  $\varepsilon = 0.03$  rad the algorithm finds 55 clusters and labels 13.5% of the events as noise. Smaller  $\varepsilon$  fragments clusters and labels more events as noise; larger  $\varepsilon$  merges boundaries and reduces noise. The largest clusters correspond exactly to known seismic zones, and the silhouette on non-noise events (3,000-event sample) is 0.326.

## Viva / Discussion Questions


**Why does DBSCAN not require  $K$ ?** Clusters are grown from density-connected core points, so their number and shape emerge from the data instead of being fixed in advance.

**What does label  $-1$  mean?** The point is neither a core nor a border point — it lies in a sparse region and is treated as noise/outlier (here: isolated earthquakes).

**Why can DBSCAN struggle with varying densities?** A single  $\varepsilon$  that is small enough for a dense cluster fragments a sparse one, while one large enough for the sparse cluster merges the dense ones.

**Why is haversine used instead of Euclidean distance?** Latitude/longitude are spherical coordinates; haversine gives true great-circle distances (radian inputs,  $\varepsilon$  in radians).

**How is  $\varepsilon$  selected?** From the knee of the  $k$ -distance graph with  $k = \text{min\_samples}$  (0.03 rad here).



**Figure 5:** DBSCAN clusters (colors) and noise (black crosses) on the global 2023 earthquake catalogue. The clusters trace the tectonically active plate boundaries.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Density-based clustering algorithms ----
from sklearn.cluster import DBSCAN, HDBSCAN
# ---- Nearest-neighbor search (used to pick epsilon) ----
from sklearn.neighbors import NearestNeighbors
# ---- Evaluation metrics ----
from sklearn.metrics import silhouette_score

# Aesthetics setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (11, 6)
plt.rcParams['figure.dpi'] = 110

np.random.seed(42)
print("Density-based libraries loaded successfully!")
```

**Listing 21:** Core numerical and plotting libraries

```
# ---- Load the real USGS earthquake catalogue (global, 2023) ----
df = pd.read_csv('../data/earthquakes.csv')
print(f"Full catalogue (magnitude >= 4.0): {len(df):,} events")
print(df.head(3).to_string(index=False))

# Keep well-recorded moderate+ events: magnitude >= 4.5
quakes = df[df['mag'] >= 4.5].reset_index(drop=True)
print(f"\nAnalysed subset (magnitude >= 4.5): {len(quakes):,} events")

# ---- Convert epicenters to radians [latitude, longitude] ----
# scikit-learn's 'haversine' metric expects radian coordinates and computes
# great-circle distance on the sphere (correct for global data).
X_rad = np.radians(quakes[['latitude', 'longitude']].values)
print("Coordinate range (deg): lat", quakes.latitude.min(), "to", quakes.latitude.max(),
      "| lon", quakes.longitude.min(), "to", quakes.longitude.max())
```

**Listing 22:** Load the real USGS earthquake catalogue (global, 2023)

```
# ---- Choosing epsilon from the k-distance graph (haversine metric) ----
# Sorted distance to each event's k-th nearest neighbor. The knee marks the
# transition from dense seismic zones to sparse, isolated events.
```

```

min_samples = 10
nbrs = NearestNeighbors(n_neighbors=min_samples, metric='haversine').fit(X_rad)
distances, _ = nbrs.kneighbors(X_rad)
k_distances = np.sort(distances[:, -1]) # radians (0.03 rad = 191 km on Earth)

plt.figure(figsize=(9, 5))
plt.plot(k_distances, color='crimson', linewidth=2.5)
plt.title(f"k-Distance Graph (k={min_samples}) for Epsilon Selection", fontsize=14, fontweight='bold')
plt.xlabel("Earthquakes Sorted by Distance", fontsize=12)
plt.ylabel(f"{min_samples}-NN Distance (radians)", fontsize=12)
plt.axhline(y=0.03, color='navy', linestyle='--', label='Selected  $\epsilon = 0.03$  rad (~191 km)')
plt.legend()
plt.tight_layout()
plt.show()

```

Listing 23: Choosing epsilon from the k-distance graph (haversine metric)

```

# ---- Fit DBSCAN with the chosen epsilon and MinPts ----
# DBSCAN grows clusters from core events, connects them via density-reachability,
# and labels everything else as noise (-1).
eps_selected = 0.03 # radians -> 0.03 * 6371 km ~ 191 km
dbscan = DBSCAN(eps=eps_selected, min_samples=min_samples, metric='haversine')
db_labels = dbscan.fit_predict(X_rad)

# Identify which events are core points (returned by the model) vs. border points
core_samples_mask = np.zeros_like(db_labels, dtype=bool)
core_samples_mask[dbscan.core_sample_indices_] = True

n_clusters_db = len(set(db_labels)) - (1 if -1 in db_labels else 0)
n_noise_db = np.sum(db_labels == -1)
print(f"DBSCAN: {n_clusters_db} clusters, {n_noise_db} noise events ({n_noise_db / len(X_rad) * 100:.1f}%)")

# ---- Visualize epicenters: clustered events in color, noise as black crosses ----
plt.figure(figsize=(12, 5.5))
palette = sns.color_palette("tab20", max(n_clusters_db, 1))
for k in range(n_clusters_db):
    mask = db_labels == k
    plt.scatter(quakes.longitude[mask], quakes.latitude[mask], s=6, color=palette[k % 20], alpha=0.75)

noise_mask = db_labels == -1
plt.scatter(quakes.longitude[noise_mask], quakes.latitude[noise_mask], s=6, color='black', marker='x', label='
    ↳ Noise (-1)')
plt.title(f"DBSCAN on Global Earthquakes 2023  $\epsilon$ ={eps_selected} rad, {n_clusters_db} clusters)", fontsize=13,
    ↳ fontweight='bold')
plt.xlabel("Longitude", fontsize=12)
plt.ylabel("Latitude", fontsize=12)
plt.legend(loc='lower left')
plt.tight_layout()
plt.show()

```

Listing 24: Fit DBSCAN with the chosen epsilon and MinPts

```

# ---- Parameter sensitivity: DBSCAN under several (eps, min_samples) settings ----
# The manual requires at least two settings so the effect of eps can be discussed.
print("eps (rad)    approx km    clusters    noise %")
for eps_test in [0.02, 0.03, 0.05]:
    lbl = DBSCAN(eps=eps_test, min_samples=10, metric='haversine').fit_predict(X_rad)
    n_clusters_test = len(set(lbl)) - (1 if -1 in lbl else 0)
    noise_pct = (lbl == -1).mean() * 100
    print(f"    {eps_test:.2f}          {eps_test * 6371:5.0f}          {n_clusters_test:3d}          {noise_pct:5.1f}%")

```

Listing 25: Parameter sensitivity: DBSCAN under several (eps, min\_samples) settings

## Experiment 6: HDBSCAN

**Dataset:** USGS earthquakes 2023 (7,638 events, magnitude  $\geq 4.5$ ; radians + haversine)

**Key parameters:** `min_cluster_size = 50`, `min_samples = 10`; no global  $\varepsilon$

**Headline result:** 40 clusters, 21.5% noise; silhouette 0.631 on non-noise events

### Objective

Apply hierarchical density-based clustering to discover clusters of different densities and identify noise without selecting a global radius.

### Suggested Dataset

The same real earthquake catalogue as Experiment 5 (7,638 events, radians + haversine metric).

### Required Libraries

scikit-learn, numpy, pandas, matplotlib.

### Brief Theory

HDBSCAN removes the need for a global  $\varepsilon$  by building a hierarchy over all density levels:

1. compute the **core distance**  $d_{\text{core}}(p)$  — the distance to a point's  $k$ -th nearest neighbor;
2. define the **mutual reachability distance**  $d_{\text{mreach}}(a, b) = \max\{d_{\text{core}}(a), d_{\text{core}}(b), d(a, b)\}$ ;
3. build a **minimum spanning tree** over the mutual-reachability graph;
4. convert the MST into a cluster hierarchy and condense it, tracking cluster births and deaths;
5. keep the clusters with the highest **stability**  $S(C) = \sum_{p \in C} (\lambda_p - \lambda_{\text{birth}}(C))$ .

The result is a flat clustering with per-point membership probabilities and noise labels, with no radius to tune.

### Procedure

1. Reuse the prepared radian coordinates.
2. Fit HDBSCAN with `min_cluster_size = 50` and `min_samples = 10`.
3. Count clusters and noise events; inspect the membership probabilities.
4. Visualize the partition.
5. Compare with DBSCAN.
6. Discuss the effect of `min_cluster_size`.

### Reference Implementation

```
hdb = HDBSCAN(min_cluster_size=50, min_samples=10, metric='haversine', copy=False)
hdb_labels = hdb.fit_predict(X_rad)
```

**Listing 26:** HDBSCAN on the earthquake data (no global epsilon)


### Evaluation / Expected Output

HDBSCAN needs no  $\varepsilon$  and handles the varying density along plate boundaries better than DBSCAN; it is more conservative (21.5% vs. 13.5% noise) because a cluster must remain stable across density scales. Its silhouette is roughly double DBSCAN's because the fixed- $\varepsilon$  algorithm chains entire arc-shaped boundaries into single elongated clusters. Membership probabilities are high ( $> 0.9$ ) deep inside seismic zones and lower near cluster edges.



**Table 10:** Density-based comparison on the 2023 earthquake catalogue.

| Algorithm                    | Clusters | Noise | Largest cluster | Silhouette   |
|------------------------------|----------|-------|-----------------|--------------|
| DBSCAN ( $\epsilon = 0.03$ ) | 55       | 13.5% | 1,629           | 0.326        |
| HDBSCAN                      | 40       | 21.5% | 560             | <b>0.631</b> |

**Figure 6:** HDBSCAN clusters and noise on the same catalogue; adaptive density splits the long DBSCAN arc clusters into compact, well-separated groups.

## Viva / Discussion Questions

**What is hierarchical about HDBSCAN?** It constructs a full hierarchy of density-based clusters across all density levels and then condenses it by tracking cluster birth and death, keeping only the most stable clusters.

**What does membership probability mean?** A per-point confidence that it belongs to its assigned cluster, derived from the point's persistence as the density threshold changes; low values indicate borderline points.

**Why is HDBSCAN useful for variable-density data?** Clusters are selected by stability across density levels rather than by one global  $\epsilon$ , so dense and sparse clusters can coexist.

**What does `min_cluster_size` control?** The smallest group that may become a cluster; increasing it reduces the number of small clusters, decreasing it fragments the map.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Density-based clustering algorithms ----
from sklearn.cluster import DBSCAN, HDBSCAN
# ---- Nearest-neighbor search (used to pick epsilon) ----
from sklearn.neighbors import NearestNeighbors
# ---- Evaluation metrics ----
from sklearn.metrics import silhouette_score

# Aesthetics setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (11, 6)
plt.rcParams['figure.dpi'] = 110

np.random.seed(42)
print("Density-based libraries loaded successfully!")
```

**Listing 27:** Core numerical and plotting libraries

```
# ---- Load the real USGS earthquake catalogue (global, 2023) ----
df = pd.read_csv('../data/earthquakes.csv')
print(f"Full catalogue (magnitude >= 4.0): {len(df):,} events")
print(df.head(3).to_string(index=False))

# Keep well-recorded moderate+ events: magnitude >= 4.5
quakes = df[df['mag'] >= 4.5].reset_index(drop=True)
print(f"\nAnalysed subset (magnitude >= 4.5): {len(quakes):,} events")

# ---- Convert epicenters to radians [latitude, longitude] ----
# scikit-learn's 'haversine' metric expects radian coordinates and computes
# great-circle distance on the sphere (correct for global data).
X_rad = np.radians(quakes[['latitude', 'longitude']].values)
print("Coordinate range (deg): lat", quakes.latitude.min(), "to", quakes.latitude.max(),
      "| lon", quakes.longitude.min(), "to", quakes.longitude.max())
```

**Listing 28:** Load the real USGS earthquake catalogue (global, 2023)

```
# ---- Fit HDBSCAN (no global epsilon required) ----
# min_cluster_size = smallest group considered a cluster; min_samples controls conservativeness.
hdb = HDBSCAN(min_cluster_size=50, min_samples=10, metric='haversine', copy=False)
hdb_labels = hdb.fit_predict(X_rad)

n_clusters_hdb = len(set(hdb_labels)) - (1 if -1 in hdb_labels else 0)
n_noise_hdb = np.sum(hdb_labels == -1)
print(f"HDBSCAN: {n_clusters_hdb} clusters, {n_noise_hdb:,} noise events ({n_noise_hdb / len(X_rad) * 100:.1f}%)")

# ---- Visualize the HDBSCAN partition ----
plt.figure(figsize=(12, 5.5))
palette = sns.color_palette("tab20", max(n_clusters_hdb, 1))
for k in range(n_clusters_hdb):
    mask = hdb_labels == k
    plt.scatter(quakes.longitude[mask], quakes.latitude[mask], s=6, color=palette[k % 20], alpha=0.75)

noise_mask = hdb_labels == -1
plt.scatter(quakes.longitude[noise_mask], quakes.latitude[noise_mask], s=6, color='black', marker='x', label='
↳ Noise (-1)')
plt.title(f"HDBSCAN on Global Earthquakes 2023 ({n_clusters_hdb} clusters, no  $\epsilon$ )", fontsize=13, fontweight='bold')
plt.xlabel("Longitude", fontsize=12)
plt.ylabel("Latitude", fontsize=12)
plt.legend(loc='lower left')
plt.tight_layout()
plt.show()
```

**Listing 29:** Fit HDBSCAN (no global epsilon required)

# **Part III**

## **Semi-Supervised Learning**

## Experiment 7: Self-training for Semi-supervised Learning

**Dataset:** Breast Cancer Wisconsin ( $569 \times 30$ , binary); training set 398 samples with only 15% labeled (61), 337 unlabeled

**Key parameters:** RBF-kernel SVC ( $C = 1$ , `probability = True`), confidence threshold  $\tau = 0.80$ , `max_iter = 15`

**Headline result:** Baseline 0.9298  $\rightarrow$  self-training 0.9415 (fully supervised ceiling 0.9766)

### Objective

Use a small labeled subset and a larger unlabeled subset, iteratively adding high-confidence pseudo-labels, and compare against a supervised baseline and the fully supervised ceiling.

### Suggested Dataset

Breast Cancer Wisconsin ( $569 \times 30$ , binary). Stratified 70/30 split; only 15% of the training labels are kept, the rest are masked as unlabeled ( $-1$ ).

### Required Libraries

scikit-learn, numpy.

### Brief Theory

Given a labeled set  $L$  and an unlabeled pool  $U$ , self-training repeats:

1. train the base classifier  $f$  on  $L$ ;
2. compute class probabilities  $f(x)$  for every  $x \in U$ ;
3. select the predictions whose maximum probability exceeds the threshold  $\tau$  and add them to  $L$  as **pseudo-labels**;
4. remove those samples from  $U$ ;
5. repeat until no new labels pass the threshold or the iteration limit is reached.

The threshold controls the purity/quantity trade-off: a low  $\tau$  admits wrong pseudo-labels (confirmation bias), while a very high  $\tau$  admits too few to help. The base estimator must expose `predict_proba`, which is why the SVC is calibrated (`probability=True`).

### Procedure

1. Split the data stratified and scale the features (fit on training data only).
2. Mask 85% of the training labels to  $-1$ .
3. Train a supervised baseline SVC on the labeled 61 samples.
4. Wrap an RBF SVC in `SelfTrainingClassifier` with  $\tau = 0.80$  and fit on the labeled plus unlabeled data; monitor the iterations.
5. Train a fully supervised SVC on all 398 labels as the ceiling.
6. Compare accuracy and F1 on the untouched test set.

### Reference Implementation

```
self_training_svc = SelfTrainingClassifier(  
    estimator=SVC(kernel='rbf', probability=True, C=1.0, random_state=42),  
    threshold=0.80, criterion='threshold', max_iter=15, verbose=True)  
self_training_svc.fit(X_train_full_scaled, y_train_semi)  # labels + unlabeled (-1)
```

**Listing 30:** Self-training with an RBF-kernel SVC

## Evaluation / Expected Output

**Table 11:** Self-training compared with the supervised baseline and the fully supervised ceiling.

| Setting                           | Accuracy      | F1            | Training samples |
|-----------------------------------|---------------|---------------|------------------|
| Supervised baseline (15% labeled) | 0.9298        | 0.9469        | 61               |
| Self-training (semi-supervised)   | <b>0.9415</b> | <b>0.9554</b> | 382              |
| Fully supervised ceiling (100%)   | 0.9766        | 0.9813        | 398              |

Pseudo-labels added per iteration: +292, +22, +5, +2, then termination `no_change` (321 new labels in total). The decaying additions indicate stable convergence rather than runaway error accumulation. Using no extra human labels, self-training improves accuracy by +1.2 points over the baseline, closing part of the gap to the fully supervised model; the remaining gap is the cost of noisy pseudo-labels.

## Viva / Discussion Questions

**What is a pseudo-label?** A predicted label for an unlabeled sample that the model accepts as a training target because its confidence exceeds the threshold.

**What is confirmation bias in self-training?** The model reinforces its own mistakes: wrong confident predictions enter the training set, make the model more confident about similar errors and can degrade performance. High thresholds and convergence monitoring mitigate this.

**Why must the test labels remain hidden during training?** Otherwise the evaluation is contaminated: test data would influence the model and the reported accuracy would no longer estimate generalization.

**What does the confidence threshold control?** The purity/quantity trade-off of pseudo-labels; 0.80 balanced the two here (convergence after four useful iterations).

**Why does the SVC need `probability=True`?** The wrapper requires `predict_proba`; a plain SVC outputs only decision values, so Platt scaling is enabled.

## Complete Code Listing

```
# ---- Suppress a known sklearn deprecation notice for SVC(probability=True) ----
import warnings
# ---- Core numerical and plotting libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Dataset + train/test utilities ----
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
# ---- Semi-supervised self-training wrapper ----
from sklearn.semi_supervised import SelfTrainingClassifier
# ---- Base estimators compared inside self-training ----
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
# ---- Evaluation metrics ----
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score,
    confusion_matrix, classification_report

# Aesthetics setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110

# The installed scikit-learn still supports SVC(probability=True) but warns; silence that specific notice
warnings.filterwarnings("ignore", message="The `probability` parameter was deprecated", category=FutureWarning)
np.random.seed(42)
print("Semi-supervised libraries imported successfully!")

# ---- Load the Breast Cancer Wisconsin dataset ----
```

```

data = load_breast_cancer()
X_full = data.data           # 30 continuous diagnostic features
y_full = data.target         # 0 = malignant, 1 = benign
feature_names = data.feature_names
target_names = data.target_names

# ---- Train / Test split (stratified to preserve class balance) ----
X_train_full, X_test, y_train_full, y_test = train_test_split(
    X_full, y_full, test_size=0.30, random_state=42, stratify=y_full
)

# ---- Standardize features (fit on train only to avoid test leakage) ----
scaler = StandardScaler()
X_train_full_scaled = scaler.fit_transform(X_train_full)
X_test_scaled = scaler.transform(X_test)

# ---- Simulate the semi-supervised setting: hide 85% of the training labels ----
unlabeled_ratio = 0.85
rng = np.random.RandomState(42)
random_unlabeled_points = rng.rand(len(y_train_full)) < unlabeled_ratio # True -> mask

y_train_semi = np.copy(y_train_full)
y_train_semi[random_unlabeled_points] = -1 # -1 is the scikit-learn convention for "unlabeled"

n_labeled = np.sum(y_train_semi != -1)
n_unlabeled = np.sum(y_train_semi == -1)

print(f"Total Training Samples: {len(y_train_full)}")
print(f" - Labeled Samples (15%): {n_labeled}")
print(f" - Unlabeled Samples (85%): {n_unlabeled}")
print(f"Holdout Test Samples: {len(y_test)}")

```

**Listing 31:** Suppress a known sklearn deprecation notice for SVC(probability=True)

```

# ---- Experiment 1: Self-Training with an RBF-kernel SVC ----

# 1. Supervised baseline trained ONLY on the ~60 labeled samples
mask_labeled = (y_train_semi != -1) # boolean mask of the labeled subset
X_labeled = X_train_full_scaled[mask_labeled]
y_labeled = y_train_semi[mask_labeled]

base_svc = SVC(kernel='rbf', probability=True, C=1.0, random_state=42)
base_svc.fit(X_labeled, y_labeled)
y_pred_baseline_svc = base_svc.predict(X_test_scaled)
acc_baseline_svc = accuracy_score(y_test, y_pred_baseline_svc)
f1_baseline_svc = f1_score(y_test, y_pred_baseline_svc)

# 2. Semi-supervised self-training: base SVC + iterative pseudo-labeling above threshold 0.80
self_training_svc = SelfTrainingClassifier(
    estimator=SVC(kernel='rbf', probability=True, C=1.0, random_state=42),
    threshold=0.80, # only accept pseudo-labels with confidence >= 0.80
    criterion='threshold', # use the confidence-threshold selection rule
    max_iter=15,
    verbose=True # print how many labels are added per iteration
)
self_training_svc.fit(X_train_full_scaled, y_train_semi)
y_pred_semi_svc = self_training_svc.predict(X_test_scaled)
acc_semi_svc = accuracy_score(y_test, y_pred_semi_svc)
f1_semi_svc = f1_score(y_test, y_pred_semi_svc)

# 3. Fully supervised "ceiling": train on 100% of the ground-truth labels
full_svc = SVC(kernel='rbf', probability=True, C=1.0, random_state=42)
full_svc.fit(X_train_full_scaled, y_train_full)
y_pred_full_svc = full_svc.predict(X_test_scaled)
acc_full_svc = accuracy_score(y_test, y_pred_full_svc)
f1_full_svc = f1_score(y_test, y_pred_full_svc)

# ---- Compare the three settings ----
print(f"\n--- SVC Evaluation Results ---")
print(f"Supervised Baseline (15% labeled): Accuracy = {acc_baseline_svc:.4f}, F1 = {f1_baseline_svc:.4f}")
print(f"Self-Training (Semi-supervised): Accuracy = {acc_semi_svc:.4f}, F1 = {f1_semi_svc:.4f}")
print(f"Fully Supervised Ceiling (100%): Accuracy = {acc_full_svc:.4f}, F1 = {f1_full_svc:.4f}")

```

**Listing 32:** Experiment 1: Self-Training with an RBF-kernel SVC

```
# ---- Inspect how self-training converged ----
# n_iter_ : number of self-training iterations performed
# termination_condition_ : why training stopped (e.g. 'no_change' = no new confident labels)
# transduction_ : final labels for every training sample (-1 = still unlabeled)
n_iter = self_training_svc.n_iter_
termination_condition = self_training_svc.termination_condition_
labeled_final_mask = (self_training_svc.transduction_ != -1)

print(f"Self-Training Convergence Summary:")
print(f" - Iterations completed: {n_iter}")
print(f" - Termination condition: {termination_condition}")
print(f" - Total samples labeled after self-training: {np.sum(labeled_final_mask)} / {len(y_train_full)}")
print(f" - Pseudo-labels newly added: {np.sum(labeled_final_mask) - n_labeled}")
```

**Listing 33:** Inspect how self-training converged

```
# ---- Compact comparison table (SVC) ----
summary_data = [
    {"Model": "SVC (Baseline 15% Labeled)", "Accuracy": acc_baseline_svc, "F1-Score": f1_baseline_svc,
     "Training Samples": n_labeled},
    {"Model": "SVC (Self-Training Semi-Supervised)", "Accuracy": acc_semi_svc, "F1-Score": f1_semi_svc,
     "Training Samples": np.sum(self_training_svc.transduction_ != -1)},
    {"Model": "SVC (Fully Supervised Ceiling 100%)", "Accuracy": acc_full_svc, "F1-Score": f1_full_svc,
     "Training Samples": len(y_train_full)},
]
summary_df = pd.DataFrame(summary_data).set_index("Model")
display(summary_df.style.highlight_max(subset=['Accuracy', 'F1-Score'], color='lightgreen'))
```

**Listing 34:** Compact comparison table (SVC)

# **Part IV**

## **Ensemble Learning**



## Experiment 8: Random Forest Regression (RFR)

**Dataset:** California Housing, 5,000-row random sample, 8 features; 80/20 split  
**Key parameters:** 150 trees, max\_depth = 12, out-of-bag scoring enabled  
**Headline result:** Test  $R^2$  0.7415 (OOB 0.7513); RMSE 0.5908; MAE 0.3999

### Objective

Predict a continuous target using an ensemble of randomized decision trees; evaluate with MAE, RMSE and  $R^2$  and inspect out-of-bag validation.

### Suggested Dataset

California Housing (scikit-learn), 5,000-row random sample, target = median house value.

### Required Libraries

scikit-learn, numpy, pandas.

### Brief Theory

Random Forest Regression trains  $B$  decision trees on bootstrap samples of the training data; at each split only a random subset of features is considered, which decorrelates the trees. The ensemble prediction is the average of the tree outputs,

$$\hat{f}_{\text{RF}}(x) = \frac{1}{B} \sum_{b=1}^B T_b(x).$$

Averaging approximately uncorrelated trees reduces variance: for pairwise correlation  $\rho$ ,  $\text{Var}(\hat{f}_{\text{RF}}) = \rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$ , so adding trees can only help. Out-of-bag (OOB) samples — the  $\approx 37\%$  not drawn in a bootstrap — provide a free internal validation estimate.

### Procedure

1. Load and sample the dataset; split into training and test sets.
2. Train `RandomForestRegressor` with 150 trees, max\_depth = 12 and OOB scoring enabled.
3. Predict the test set and compute  $R^2$ , RMSE and MAE.
4. Report the OOB  $R^2$  and compare it with the test  $R^2$  to check overfitting.
5. Inspect feature importances.

### Reference Implementation

```
rf_reg = RandomForestRegressor(n_estimators=150, max_depth=12,
                              oob_score=True, random_state=42, n_jobs=-1)
rf_reg.fit(X_train_reg, y_train_reg)
print(rf_reg.oob_score_, r2_score(y_test_reg, rf_reg.predict(X_test_reg)))
```

**Listing 35:** Random Forest Regression with out-of-bag validation

### Evaluation / Expected Output

The OOB  $R^2$  (0.7513) is very close to the test  $R^2$  (0.7415), so the model is not overfitting and the OOB estimate is a reliable free validation. Feature importances are dominated by median income and location (latitude/longitude), consistent with housing economics.

**Table 12:** Random Forest Regression on the California Housing sample.

| Metric     | Value                                    |
|------------|------------------------------------------|
| OOB $R^2$  | 0.7513                                   |
| Test $R^2$ | 0.7415                                   |
| Test RMSE  | 0.5908 (in \$100k units $\approx$ \$59k) |
| Test MAE   | 0.3999 (in \$100k units $\approx$ \$40k) |

## Viva / Discussion Questions

**Why does a forest reduce variance?** Averaging  $B$  approximately uncorrelated trees gives  $\rho\sigma^2 + (1 - \rho)\sigma^2/B$ ; as  $B$  grows the second term vanishes, so predictions are more stable than a single tree's.

**What is bagging?** Bootstrap aggregating: each model is trained on a bootstrap sample of the data and the predictions are combined (average for regression, vote for classification).

**Why can feature importance be misleading?** Impurity-based importance is biased toward high-cardinality/continuous features, is computed on training data (it implies no causality), and correlated features share or dilute importance arbitrarily.

## Complete Code Listing

```
# ---- Core libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier, AdaBoostRegressor, AdaBoostClassifier
from sklearn.metrics import r2_score, root_mean_squared_error, mean_absolute_error, accuracy_score, f1_score,
    ↪ roc_auc_score

import xgboost as xgb
import catboost as cb

# Styling setup
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110
np.random.seed(42)
```

**Listing 36:** Core libraries

```
# ---- Load the California Housing dataset (real-world regression benchmark) ----
housing = fetch_california_housing(as_frame=True)
# Sample 5,000 rows so every ensemble trains quickly while remaining representative
df_reg = housing.frame.sample(n=5000, random_state=42)

X_reg = df_reg[housing.feature_names] # 8 socio-economic / geographic features
y_reg = df_reg['MedHouseVal']         # target: median house value

# ---- Train/test split (80% / 20%) ----
X_train_reg, X_test_reg, y_train_reg, y_test_reg = train_test_split(
    X_reg, y_reg, test_size=0.20, random_state=42
)

print(f"Regression Training Samples: {X_train_reg.shape[0]}, Test Samples: {X_test_reg.shape[0]}")
df_reg.head()
```

**Listing 37:** Load the California Housing dataset (real-world regression benchmark)

```
# ---- Random Forest Regression (bagging of decision trees) ----
# n_estimators = number of trees, oob_score = estimate accuracy from out-of-bag samples
rf_reg = RandomForestRegressor(n_estimators=150, max_depth=12, oob_score=True, random_state=42, n_jobs=-1)
rf_reg.fit(X_train_reg, y_train_reg)
```

```
# ---- Predict and evaluate ----
y_pred_rf = rf_reg.predict(X_test_reg)
r2_rf = r2_score(y_test_reg, y_pred_rf)           # explained variance (1.0 = perfect)
rmse_rf = root_mean_squared_error(y_test_reg, y_pred_rf) # error in target units
mae_rf = mean_absolute_error(y_test_reg, y_pred_rf)    # robust average error

print(f"Random Forest Regressor Results:")
print(f" - Out-of-Bag (OOB) R2 Score: {rf_reg.oob_score_:.4f}")
print(f" - Test R2 Score:           {r2_rf:.4f}")
print(f" - Test RMSE:               {rmse_rf:.4f}")
print(f" - Test MAE:                {mae_rf:.4f}")
```

**Listing 38:** Random Forest Regression (bagging of decision trees)

## Experiment 9: Random Forest Classification (RFC)

**Dataset:** Heart Disease (`data/heart_disease.csv`, 303 rows, 13 features); stratified 75/25 split  
**Key parameters:** 120 trees, `max_depth = 8`; one-hot encoded categorical features  
**Headline result:** Accuracy 0.7763; F1 0.8046; ROC-AUC 0.8704

### Objective

Classify observations using an ensemble of randomized decision trees, with a stratified split and classification metrics.

### Suggested Dataset

Heart Disease (`data/heart_disease.csv`), 303 rows, 13 features plus the binary target.

### Required Libraries

scikit-learn, numpy, pandas.

### Brief Theory

Random Forest Classification aggregates tree votes. Each tree votes for one class and the ensemble predicts the most frequent class,

$$\hat{C}_{\text{RF}}(x) = \arg \max_k \sum_{b=1}^B \mathbb{I}(T_b(x) = k).$$

Bootstrap sampling plus feature subsampling decorrelates the trees, reducing variance relative to a single decision tree. The fraction of trees voting for the positive class also provides a class probability, which is used for ROC-AUC.

### Procedure

1. Load the heart data; one-hot encode the categorical features.
2. Use a stratified 75/25 split to preserve the disease ratio.
3. Train `RandomForestClassifier` with 120 trees and `max_depth = 8`.
4. Predict labels and probabilities on the test set.
5. Compute accuracy, F1 and ROC-AUC.

### Reference Implementation

```
rfc = RandomForestClassifier(n_estimators=120, max_depth=8, random_state=42, n_jobs=-1)
rfc.fit(X_train_clf, y_train_clf)
y_prob_rfc = rfc.predict_proba(X_test_clf)[:, 1]
```

**Listing 39:** Random Forest Classification with stratified evaluation

### Evaluation / Expected Output

**Table 13:** Random Forest Classification on Heart Disease (76 test samples).

| Metric   | Value  |
|----------|--------|
| Accuracy | 0.7763 |
| F1-score | 0.8046 |
| ROC-AUC  | 0.8704 |

ROC-AUC 0.87 indicates good ranking quality on a small dataset; accuracy 0.776 is typical for heart-disease prediction with these features and reflects the problem's noise rather than a modelling failure. F1 and AUC are safer summaries than accuracy alone because the classes are not perfectly balanced.

## Viva / Discussion Questions

**RFR vs RFC?** Regression averages the trees' numeric outputs; classification takes a majority vote (and averages class probabilities for scoring).

**What is majority voting?** Each tree votes for one class and the most frequent class becomes the ensemble prediction; soft voting averages the probabilities instead.

**How can class imbalance affect accuracy?** With a rare positive class, a model can predict the majority class everywhere and still show high accuracy; precision, recall, F1 and AUC expose the failure.

## Complete Code Listing

```
# ---- Core libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier, AdaBoostRegressor, AdaBoostClassifier
from sklearn.metrics import r2_score, root_mean_squared_error, mean_absolute_error, accuracy_score, f1_score,
    ↪ roc_auc_score

import xgboost as xgb
import catboost as cb

# Styling setup
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110
np.random.seed(42)
```

**Listing 40:** Core libraries

```
# ---- Load the Heart Disease dataset (real-world classification benchmark) ----
df_clf = pd.read_csv('../data/heart_disease.csv')
print(f"Heart Disease dataset: {df_clf.shape[0]} rows x {df_clf.shape[1]} columns")

# Separate target and features (handle either 'target' or 'Target' as the column name)
target_col = 'target' if 'target' in df_clf.columns else 'Target'
y_clf = df_clf[target_col].values
X_clf_raw = df_clf.drop(columns=[target_col])

# One-hot encode categorical variables for RFC / XGBoost / AdaBoost
X_clf_encoded = pd.get_dummies(X_clf_raw, drop_first=True)

# Stratified split keeps the disease/no-disease ratio identical in train and test
X_train_clf, X_test_clf, y_train_clf, y_test_clf = train_test_split(
    X_clf_encoded, y_clf, test_size=0.25, random_state=42, stratify=y_clf
)

print(f"Classification Training Set: {X_train_clf.shape}, Test Set: {X_test_clf.shape}")

# ---- Random Forest Classification (majority vote of trees) ----
rfc = RandomForestClassifier(n_estimators=120, max_depth=8, random_state=42, n_jobs=-1)
rfc.fit(X_train_clf, y_train_clf)

# Predictions and probability of the positive (disease) class
y_pred_rfc = rfc.predict(X_test_clf)
y_prob_rfc = rfc.predict_proba(X_test_clf)[:, 1]

# ---- Evaluate: accuracy, F1 and ROC-AUC ----
acc_rfc = accuracy_score(y_test_clf, y_pred_rfc)
f1_rfc = f1_score(y_test_clf, y_pred_rfc)
```

```
auc_rfc = roc_auc_score(y_test_clf, y_prob_rfc)

print(f"Random Forest Classifier: Accuracy = {acc_rfc:.4f}, F1 = {f1_rfc:.4f}, ROC-AUC = {auc_rfc:.4f}")
```

**Listing 41:** Load the Heart Disease dataset (real-world classification benchmark)

## Experiment 10: XGBoost

**Dataset:** Heart Disease (same stratified 75/25 split as RFC)

**Key parameters:** 120 trees, learning rate 0.08, max\_depth = 4, subsample / colsample = 0.8

**Headline result:** Accuracy 0.7763; F1 0.8132; ROC-AUC 0.8509

### Objective

Train a gradient-boosted tree classifier and analyze the effect of boosting hyperparameters.

### Suggested Dataset

Heart Disease (data/heart\_disease.csv), the same stratified split used for the other ensembles.

### Required Libraries

xgboost, scikit-learn, numpy, pandas.

### Brief Theory

XGBoost builds trees sequentially, each fitting a second-order Taylor expansion of the loss:

$$\mathcal{L}^{(t)} \approx \sum_{i=1}^n \left[ g_i f_t(x_i) + \frac{1}{2} h_i f_t^2(x_i) \right] + \Omega(f_t), \quad \Omega(f) = \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2,$$

where  $g_i$  and  $h_i$  are the first and second derivatives of the loss with respect to the current prediction. The regularization term  $\Omega$  penalizes tree complexity and leaf weights, while row/column subsampling adds stochasticity; both control overfitting.

### Procedure

1. Prepare the one-hot encoded heart data and the same stratified split.
2. Fit `XGBClassifier` with 120 trees, learning rate 0.08, depth 4, subsample and column sample 0.8, and the logistic-loss objective.
3. Predict labels and probabilities on the test set.
4. Report accuracy, F1 and ROC-AUC.
5. Discuss the complexity/generalization trade-off.

### Reference Implementation

```
xgb_clf = xgb.XGBClassifier(n_estimators=120, learning_rate=0.08, max_depth=4,
                           subsample=0.8, colsample_bytree=0.8,
                           random_state=42, eval_metric='logloss')
xgb_clf.fit(X_train_clf, y_train_clf)
y_prob_xgb_c = xgb_clf.predict_proba(X_test_clf)[:, 1]
```

**Listing 42:** XGBoost classification on the heart data

### Evaluation / Expected Output

XGBoost matches RFC on accuracy with a slightly higher F1 but a slightly lower AUC on this small test set; differences of 0.01–0.02 are within noise for 76 samples. On the larger regression benchmark, gradient boosting clearly leads (Experiment 8 comparison), which is the usual pattern for tabular data.

**Table 14:** XGBoost classification on Heart Disease (76 test samples).

| Metric   | Value  |
|----------|--------|
| Accuracy | 0.7763 |
| F1-score | 0.8132 |
| ROC-AUC  | 0.8509 |

## Viva / Discussion Questions

**Bagging vs boosting?** Bagging trains models in parallel on bootstrap samples and combines them to reduce variance; boosting trains models sequentially so each corrects the previous ensemble's errors, mainly reducing bias.

**What does the learning rate do?** It shrinks each tree's contribution; smaller values need more trees but usually generalize better, larger values fit faster and can overfit.

**Why can boosting overfit?** Later trees keep fitting remaining errors, including label noise, so the ensemble can become too complex; depth limits, regularization, subsampling and early stopping control this.

## Complete Code Listing

```
# ---- Core libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier, AdaBoostRegressor, AdaBoostClassifier
from sklearn.metrics import r2_score, root_mean_squared_error, mean_absolute_error, accuracy_score, f1_score,
    ↪ roc_auc_score

import xgboost as xgb
import catboost as cb

# Styling setup
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110
np.random.seed(42)
```

**Listing 43:** Core libraries

```
# ---- Load the Heart Disease dataset (real-world classification benchmark) ----
df_clf = pd.read_csv('../data/heart_disease.csv')
print(f"Heart Disease dataset: {df_clf.shape[0]} rows x {df_clf.shape[1]} columns")

# Separate target and features (handle either 'target' or 'Target' as the column name)
target_col = 'target' if 'target' in df_clf.columns else 'Target'
y_clf = df_clf[target_col].values
X_clf_raw = df_clf.drop(columns=[target_col])

# One-hot encode categorical variables for RFC / XGBoost / AdaBoost
X_clf_encoded = pd.get_dummies(X_clf_raw, drop_first=True)

# Stratified split keeps the disease/no-disease ratio identical in train and test
X_train_clf, X_test_clf, y_train_clf, y_test_clf = train_test_split(
    X_clf_encoded, y_clf, test_size=0.25, random_state=42, stratify=y_clf
)

print(f"Classification Training Set: {X_train_clf.shape}, Test Set: {X_test_clf.shape}")

# ---- Random Forest Classification (majority vote of trees) ----
rfc = RandomForestClassifier(n_estimators=120, max_depth=8, random_state=42, n_jobs=-1)
rfc.fit(X_train_clf, y_train_clf)

# Predictions and probability of the positive (disease) class
y_pred_rfc = rfc.predict(X_test_clf)
```



```
y_prob_rfc = rfc.predict_proba(X_test_clf)[:, 1]

# ---- Evaluate: accuracy, F1 and ROC-AUC ----
acc_rfc = accuracy_score(y_test_clf, y_pred_rfc)
f1_rfc = f1_score(y_test_clf, y_pred_rfc)
auc_rfc = roc_auc_score(y_test_clf, y_prob_rfc)

print(f"Random Forest Classifier: Accuracy = {acc_rfc:.4f}, F1 = {f1_rfc:.4f}, ROC-AUC = {auc_rfc:.4f}")
```

**Listing 44:** Load the Heart Disease dataset (real-world classification benchmark)

```
# ---- XGBoost classification on the same stratified split ----
xgb_clf = xgb.XGBClassifier(
    n_estimators=120, learning_rate=0.08, max_depth=4,
    subsample=0.8, colsample_bytree=0.8,
    random_state=42, eval_metric='logloss'
)
xgb_clf.fit(X_train_clf, y_train_clf)

y_pred_xgb_c = xgb_clf.predict(X_test_clf)
y_prob_xgb_c = xgb_clf.predict_proba(X_test_clf)[:, 1]

acc_xgb_c = accuracy_score(y_test_clf, y_pred_xgb_c)
f1_xgb_c = f1_score(y_test_clf, y_pred_xgb_c)
auc_xgb_c = roc_auc_score(y_test_clf, y_prob_xgb_c)
print(f"XGBoost Classifier: Accuracy = {acc_xgb_c:.4f}, F1 = {f1_xgb_c:.4f}, ROC-AUC = {auc_xgb_c:.4f}")
```

**Listing 45:** XGBoost classification on the same stratified split

## Experiment 11: AdaBoost

**Dataset:** Heart Disease (same stratified 75/25 split)

**Key parameters:** 100 decision stumps, learning rate 0.1

**Headline result:** Accuracy 0.7895; F1 0.8182; ROC-AUC 0.8714 (best of the four ensembles)

### Objective

Build an adaptive boosting classifier and study how weak learners combine into a stronger ensemble.

### Suggested Dataset

Heart Disease (`data/heart_disease.csv`), same split as the other ensembles.

### Required Libraries

scikit-learn, numpy, pandas.

### Brief Theory

AdaBoost trains weak learners (decision stumps) sequentially on re-weighted data. Misclassified samples receive larger weights, so each new learner focuses on the current errors. At round  $m$  the weighted error  $\epsilon_m$  determines the learner weight

$$\alpha_m = \frac{1}{2} \ln \left( \frac{1 - \epsilon_m}{\epsilon_m} \right),$$

and the sample weights are updated as  $w_i^{(m+1)} = w_i^{(m)} \exp(-\alpha_m y_i h_m(x_i))$ . The final prediction is a weighted vote of all stumps. Because the exponential loss grows quickly for hard points, AdaBoost is sensitive to label noise and outliers.

### Procedure

1. Prepare the heart data and the stratified split.
2. Fit `AdaBoostClassifier` with 100 stumps and learning rate 0.1.
3. Predict labels and probabilities on the test set.
4. Report accuracy, F1 and ROC-AUC; compare with RFC and XGBoost.

### Reference Implementation

```
ada_clf = AdaBoostClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
ada_clf.fit(X_train_clf, y_train_clf)
y_prob_ada_c = ada_clf.predict_proba(X_test_clf)[:, 1]
```

**Listing 46:** AdaBoost classification

### Evaluation / Expected Output

**Table 15:** AdaBoost classification on Heart Disease (76 test samples).

| Metric   | Value         |
|----------|---------------|
| Accuracy | <b>0.7895</b> |
| F1-score | <b>0.8182</b> |
| ROC-AUC  | <b>0.8714</b> |

AdaBoost with shallow stumps slightly outperforms the deeper models on this small dataset — a common outcome

when the signal is mostly additive and the sample size is limited. It remains a cheap, strong baseline; its weakness is sensitivity to label noise, where regularized gradient boosting is more robust.

## Viva / Discussion Questions

**Why use weak learners?** Individually they are simple and low-variance; boosting combines many of them so the ensemble can fit complex boundaries while each step stays simple.

**How does AdaBoost focus on difficult examples?** After each round, weights of misclassified samples increase, so the next learner optimizes a distribution that emphasizes current errors.

**What is the effect of noisy labels?** Mislabeled points look “hard”, receive ever-larger weights and can dominate training, degrading accuracy — AdaBoost is known to be noise-sensitive.

## Complete Code Listing

```
# ---- Core libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier, AdaBoostRegressor, AdaBoostClassifier
from sklearn.metrics import r2_score, root_mean_squared_error, mean_absolute_error, accuracy_score, f1_score,
    ↪ roc_auc_score

import xgboost as xgb
import catboost as cb

# Styling setup
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110
np.random.seed(42)
```

**Listing 47:** Core libraries

```
# ---- Load the Heart Disease dataset (real-world classification benchmark) ----
df_clf = pd.read_csv('../data/heart_disease.csv')
print(f"Heart Disease dataset: {df_clf.shape[0]} rows x {df_clf.shape[1]} columns")

# Separate target and features (handle either 'target' or 'Target' as the column name)
target_col = 'target' if 'target' in df_clf.columns else 'Target'
y_clf = df_clf[target_col].values
X_clf_raw = df_clf.drop(columns=[target_col])

# One-hot encode categorical variables for RFC / XGBoost / AdaBoost
X_clf_encoded = pd.get_dummies(X_clf_raw, drop_first=True)

# Stratified split keeps the disease/no-disease ratio identical in train and test
X_train_clf, X_test_clf, y_train_clf, y_test_clf = train_test_split(
    X_clf_encoded, y_clf, test_size=0.25, random_state=42, stratify=y_clf
)

print(f"Classification Training Set: {X_train_clf.shape}, Test Set: {X_test_clf.shape}")

# ---- Random Forest Classification (majority vote of trees) ----
rfc = RandomForestClassifier(n_estimators=120, max_depth=8, random_state=42, n_jobs=-1)
rfc.fit(X_train_clf, y_train_clf)

# Predictions and probability of the positive (disease) class
y_pred_rfc = rfc.predict(X_test_clf)
y_prob_rfc = rfc.predict_proba(X_test_clf)[:, 1]

# ---- Evaluate: accuracy, F1 and ROC-AUC ----
acc_rfc = accuracy_score(y_test_clf, y_pred_rfc)
f1_rfc = f1_score(y_test_clf, y_pred_rfc)
auc_rfc = roc_auc_score(y_test_clf, y_prob_rfc)
```

```
print(f"Random Forest Classifier: Accuracy = {acc_rfc:.4f}, F1 = {f1_rfc:.4f}, ROC-AUC = {auc_rfc:.4f}")
```

**Listing 48:** Load the Heart Disease dataset (real-world classification benchmark)

```
# ---- AdaBoost classification ----
ada_clf = AdaBoostClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
ada_clf.fit(X_train_clf, y_train_clf)

y_pred_ada_c = ada_clf.predict(X_test_clf)
y_prob_ada_c = ada_clf.predict_proba(X_test_clf)[:, 1]

acc_ada_c = accuracy_score(y_test_clf, y_pred_ada_c)
f1_ada_c = f1_score(y_test_clf, y_pred_ada_c)
auc_ada_c = roc_auc_score(y_test_clf, y_prob_ada_c)
print(f"AdaBoost Classifier: Accuracy = {acc_ada_c:.4f}, F1 = {f1_ada_c:.4f}, ROC-AUC = {auc_ada_c:.4f}")
```

**Listing 49:** AdaBoost classification

## Experiment 12: CatBoost

**Dataset:** Heart Disease (raw table with categorical features: sex, cp, fbs, restecg, exang, slope, ca, thal)  
**Key parameters:** 150 iterations, depth 5, learning rate 0.08, native `cat_features`  
**Headline result:** Accuracy 0.7632; F1 0.7955; ROC-AUC 0.8551

### Objective

Train a gradient-boosted decision-tree model with native handling of categorical variables.

### Suggested Dataset

Heart Disease (`data/heart_disease.csv`) using the raw (unencoded) table so CatBoost can process categorical columns directly.

### Required Libraries

catboost, pandas, scikit-learn.

### Brief Theory

CatBoost is a gradient-boosting library designed for tables that mix numerical and categorical features. It uses **ordered boosting** — leaf values are computed on random permutations of the data — to avoid target leakage and prediction shift, and builds **oblivious (symmetric) trees**, where the same split condition is used at every node of a depth level. Categorical features are encoded internally with ordered target statistics, so the categorical columns must be declared through `cat_features` and the target must never leak into the encoding.

### Procedure

1. Identify the categorical columns in the raw heart table and cast them to string.
2. Use the same stratified 75/25 split on the raw features.
3. Fit `CatBoostClassifier` with 150 iterations, depth 5, learning rate 0.08 and the categorical columns declared.
4. Predict labels and probabilities; report accuracy, F1 and ROC-AUC.
5. Compare all four classifiers on the same split.

### Reference Implementation

```
categorical_cols = X_clf_raw.select_dtypes(include=['object', 'category']).columns.tolist()
cb_clf = cb.CatBoostClassifier(iterations=150, learning_rate=0.08, depth=5,
                              cat_features=categorical_cols, random_seed=42, verbose=False)
cb_clf.fit(X_tr_cb, y_tr_cb)
```

**Listing 50:** CatBoost with native categorical features

### Evaluation / Expected Output

**Table 16:** Ensemble classification comparison on the same Heart Disease split (76 test samples).

| Model               | Accuracy      | F1            | ROC-AUC       |
|---------------------|---------------|---------------|---------------|
| Random Forest (RFC) | 0.7763        | 0.8046        | 0.8704        |
| XGBoost             | 0.7763        | 0.8132        | 0.8509        |
| AdaBoost            | <b>0.7895</b> | <b>0.8182</b> | <b>0.8714</b> |
| CatBoost            | 0.7632        | 0.7955        | 0.8551        |

CatBoost performs at a similar level to the other ensembles without any manual one-hot encoding — its native categorical handling does the encoding internally. On this tiny dataset the advanced target statistics cannot show their full advantage (they matter most with high-cardinality categories, e.g. thousands of city or user IDs); all differences are within noise for 76 test samples.

## Viva / Discussion Questions

**How does CatBoost treat categorical features?** It computes ordered target statistics (the mean target per category on random permutations, with prior smoothing) rather than one-hot vectors.

**Why should categories not be encoded using target leakage?** Plain target encoding uses each row's own target when encoding that row, which leaks label information and inflates training performance; ordered statistics avoid this.

**CatBoost vs ordinary one-hot encoding?** One-hot explodes dimensionality for high-cardinality features and loses category statistics; CatBoost encodes compactly and exploits target–category relationships with leakage control.

## Complete Code Listing

```
# ---- Core libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier, AdaBoostRegressor, AdaBoostClassifier
from sklearn.metrics import r2_score, root_mean_squared_error, mean_absolute_error, accuracy_score, f1_score,
    ↪ roc_auc_score

import xgboost as xgb
import catboost as cb

# Styling setup
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110
np.random.seed(42)
```

**Listing 51:** Core libraries

```
# ---- Load the Heart Disease dataset (real-world classification benchmark) ----
df_clf = pd.read_csv('../data/heart_disease.csv')
print(f"Heart Disease dataset: {df_clf.shape[0]} rows x {df_clf.shape[1]} columns")

# Separate target and features (handle either 'target' or 'Target' as the column name)
target_col = 'target' if 'target' in df_clf.columns else 'Target'
y_clf = df_clf[target_col].values
X_clf_raw = df_clf.drop(columns=[target_col])

# One-hot encode categorical variables for RFC / XGBoost / AdaBoost
X_clf_encoded = pd.get_dummies(X_clf_raw, drop_first=True)

# Stratified split keeps the disease/no-disease ratio identical in train and test
X_train_clf, X_test_clf, y_train_clf, y_test_clf = train_test_split(
    X_clf_encoded, y_clf, test_size=0.25, random_state=42, stratify=y_clf
)

print(f"Classification Training Set: {X_train_clf.shape}, Test Set: {X_test_clf.shape}")

# ---- Random Forest Classification (majority vote of trees) ----
rfc = RandomForestClassifier(n_estimators=120, max_depth=8, random_state=42, n_jobs=-1)
rfc.fit(X_train_clf, y_train_clf)

# Predictions and probability of the positive (disease) class
y_pred_rfc = rfc.predict(X_test_clf)
y_prob_rfc = rfc.predict_proba(X_test_clf)[:, 1]

# ---- Evaluate: accuracy, F1 and ROC-AUC ----
acc_rfc = accuracy_score(y_test_clf, y_pred_rfc)
```

```
fl_rfc = fl_score(y_test_clf, y_pred_rfc)
auc_rfc = roc_auc_score(y_test_clf, y_prob_rfc)

print(f"Random Forest Classifier: Accuracy = {acc_rfc:.4f}, F1 = {fl_rfc:.4f}, ROC-AUC = {auc_rfc:.4f}")
```

**Listing 52:** Load the Heart Disease dataset (real-world classification benchmark)

```
# ---- CatBoost classification with native categorical features ----
# Identify categorical columns in the raw (unencoded) heart data
categorical_cols = X_clf_raw.select_dtypes(include=['object', 'category']).columns.tolist()
X_raw_cb = X_clf_raw.copy()
for col in categorical_cols:
    X_raw_cb[col] = X_raw_cb[col].astype(str)

# Same stratified split on the raw feature table
X_tr_cb, X_te_cb, y_tr_cb, y_te_cb = train_test_split(
    X_raw_cb, y_clf, test_size=0.25, random_state=42, stratify=y_clf
)

cb_clf = cb.CatBoostClassifier(
    iterations=150, learning_rate=0.08, depth=5,
    cat_features=categorical_cols, random_seed=42, verbose=False
)
cb_clf.fit(X_tr_cb, y_tr_cb)

y_pred_cb_c = cb_clf.predict(X_te_cb)
y_prob_cb_c = cb_clf.predict_proba(X_te_cb)[:, 1]

acc_cb_c = accuracy_score(y_te_cb, y_pred_cb_c)
fl_cb_c = fl_score(y_te_cb, y_pred_cb_c)
auc_cb_c = roc_auc_score(y_te_cb, y_prob_cb_c)
print(f"CatBoost Classifier: Accuracy = {acc_cb_c:.4f}, F1 = {fl_cb_c:.4f}, ROC-AUC = {auc_cb_c:.4f}")
```

**Listing 53:** CatBoost classification with native categorical features

## **Part V**

# **Neural Networks: Multilayer Perceptron**



## Experiment 13: Multilayer Perceptron (MLP)

**Dataset:** Digits (scikit-learn): 1,797 samples,  $8 \times 8$  images (64 features), 10 classes; 70/15/15 split

**Key parameters:** scikit-learn (128, 64) ReLU + Adam + early stopping; PyTorch 64–128–64–10 with BatchNorm and Dropout

**Headline result:** scikit-learn 0.9630; PyTorch 0.9815 test accuracy (GPU)

### Objective

Implement feed-forward neural networks for handwritten-digit classification and study hidden-layer size and regularization.

### Suggested Dataset

Digits (scikit-learn): 1,797 samples,  $8 \times 8$  grayscale images (64 features), 10 classes; stratified 70/15/15 train/validation/test split, standardized on the training data only.

### Required Libraries

scikit-learn, PyTorch, numpy, pandas, matplotlib.

### Brief Theory

An MLP stacks affine transformations with non-linear activations,

$$a^{(l)} = \sigma(W^{(l)}a^{(l-1)} + b^{(l)}), \quad \hat{y} = \text{softmax}(W^{(L+1)}a^{(L)} + b^{(L+1)}),$$

and is trained by backpropagation to minimize the categorical cross-entropy

$$\mathcal{L} = - \sum_{i=1}^N \sum_{k=1}^K y_{ik} \ln \hat{y}_{ik} + \frac{\alpha}{2} \sum_l \|W^{(l)}\|_F^2.$$

**Batch normalization** stabilizes layer inputs and speeds training, **dropout** randomly zeroes neurons to regularize, and **early stopping** / learning-rate scheduling stop training when validation performance stalls.

### Procedure

1. Load the digits, split 70/15/15 and standardize (fit on training data only).
2. Train a scikit-learn `MLPClassifier` with hidden layers (128, 64), ReLU, Adam,  $L_2$  regularization and early stopping.
3. Build a custom PyTorch MLP: 64–128–64–10 with BatchNorm and Dropout, cross-entropy loss, Adam with weight decay, 50 epochs and a plateau scheduler.
4. Train with the validation loop and plot the loss curve.
5. Evaluate the test set and report the per-class classification report.

### Reference Implementation

```
mlp_sklearn = MLPClassifier(hidden_layer_sizes=(128, 64), activation='relu', solver='adam',
                             alpha=0.001, batch_size=64, learning_rate_init=0.005,
                             max_iter=150, early_stopping=True, validation_fraction=0.15)


self.net = nn.Sequential(
    nn.Linear(64, 128), nn.BatchNorm1d(128), nn.ReLU(), nn.Dropout(0.25),
    nn.Linear(128, 64), nn.BatchNorm1d(64), nn.ReLU(), nn.Dropout(0.20),
    nn.Linear(64, 10))
```

**Listing 54:** The two MLP architectures

## Evaluation / Expected Output

**Table 17:** MLP implementations on the digits test set (270 images).

| Model                     | Test accuracy | Notes                                             |
|---------------------------|---------------|---------------------------------------------------|
| Scikit-learn MLP (128–64) | 0.9630        | early stopping after 15 iterations                |
| PyTorch deep MLP          | <b>0.9815</b> | GPU (CUDA), 50 epochs, validation accuracy 0.9889 |



**Figure 7:** Training loss curve of the scikit-learn MLP; early stopping halted training after 15 iterations.

Training loss fell from 1.20 to about 0.02 while the validation loss tracked it closely, so there is no severe overfitting. Per-class precision and recall are near 1.00 for almost all digits; the residual errors are visually similar digits. The custom PyTorch model beats the scikit-learn baseline by about two points thanks to BatchNorm, Dropout and scheduled training over more epochs.

## Viva / Discussion Questions

**Why is scaling important for MLP?** Gradient descent converges poorly when input features have very different scales; standardization equalizes their influence and stabilizes the loss surface.

**What is backpropagation?** Reverse-mode application of the chain rule that computes the gradient of the loss with respect to every weight, layer by layer.

**What causes vanishing/exploding gradients?** Repeated multiplication of small (or large) Jacobians through many layers; ReLU, BatchNorm, careful initialization and residual connections mitigate it.

**Why is the PyTorch model better here?** BatchNorm and Dropout plus a longer, scheduled training; the scikit-learn model stops early at iteration 15.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Dataset + preprocessing/metrics ----
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
# ---- Scikit-Learn MLP implementation ----
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```

# ---- PyTorch for the custom deep MLP ----
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

# Styling setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110

# Fix seeds for reproducible training
torch.manual_seed(42)
torch.cuda.manual_seed_all(42) # reproducible GPU training
np.random.seed(42)
print(f"PyTorch version: {torch.__version__}")
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu') # use GPU if available
print(f"Computation Device: {device}")

# ---- Load the handwritten digits dataset (8x8 grayscale images) ----
digits = load_digits()
X_raw = digits.data # flattened images: 1797 samples x 64 pixel features
y = digits.target # digit class 0..9

print(f"Dataset Shape: {X_raw.shape}, Target Classes: {np.unique(y)}")
# ---- Train / Validation / Test split (70% / 15% / 15%) ----
# First cut off the test set...
X_train_val, X_test, y_train_val, y_test = train_test_split(
    X_raw, y, test_size=0.15, random_state=42, stratify=y
)
# ...then split the remainder into train and validation
X_train, X_val, y_train, y_val = train_test_split(
    X_train_val, y_train_val, test_size=0.1765, random_state=42, stratify=y_train_val
)

# ---- Feature scaling: fit on training data only (avoid validation/test leakage) ----
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_val_scaled = scaler.transform(X_val)
X_test_scaled = scaler.transform(X_test)

print(f"Training samples: {X_train_scaled.shape[0]}")
print(f"Validation samples: {X_val_scaled.shape[0]}")
print(f"Test samples: {X_test_scaled.shape[0]}")

```

Listing 55: Core numerical and plotting libraries

```

# ---- Scikit-Learn MLPClassifier: 64 -> 128 -> 64 -> 10 ----
mlp_sklearn = MLPClassifier(
    hidden_layer_sizes=(128, 64), # two hidden layers
    activation='relu',
    solver='adam', # adaptive gradient optimizer
    alpha=0.001, # L2 regularization strength
    batch_size=64,
    learning_rate_init=0.005,
    max_iter=150,
    early_stopping=True, # stop when validation score stops improving
    validation_fraction=0.15, # internal validation split for early stopping
    random_state=42
)

# Train (forward pass + backpropagation happen inside .fit)
mlp_sklearn.fit(X_train_scaled, y_train)

# Evaluate on the untouched test set
y_pred_sk = mlp_sklearn.predict(X_test_scaled)
acc_sk = accuracy_score(y_test, y_pred_sk)

print(f"Scikit-Learn MLP Test Accuracy: {acc_sk:.4f}")
print(f"Converged after {mlp_sklearn.n_iter_} iterations")

# ---- Plot the training loss curve recorded by scikit-learn ----
plt.figure(figsize=(8, 4.5))

```

```
plt.plot(mlp_sklearn.loss_curve_, color='crimson', lw=2.2, label='Training Loss')
if hasattr(mlp_sklearn, 'validation_scores_'):
    # validation_scores_ stores accuracy; convert to error for a comparable scale
    plt.plot([1 - s for s in mlp_sklearn.validation_scores_], color='navy', lw=2.2, linestyle='--', label='
    ↳ Validation Error')
plt.title("Scikit-Learn MLP Training Loss Curve", fontsize=13, fontweight='bold')
plt.xlabel("Epochs / Iterations", fontsize=11)
plt.ylabel("Loss", fontsize=11)
plt.legend()
plt.tight_layout()
plt.show()
```

Listing 56: Scikit-Learn MLPClassifier: 64 -&gt; 128 -&gt; 64 -&gt; 10

```
# ---- Convert scaled NumPy arrays into PyTorch tensors + DataLoaders ----
train_dataset = TensorDataset(torch.tensor(X_train_scaled, dtype=torch.float32), torch.tensor(y_train, dtype=torch.
    ↳ long))
val_dataset = TensorDataset(torch.tensor(X_val_scaled, dtype=torch.float32), torch.tensor(y_val, dtype=torch.
    ↳ long))
test_dataset = TensorDataset(torch.tensor(X_test_scaled, dtype=torch.float32), torch.tensor(y_test, dtype=torch.
    ↳ long))

# DataLoaders handle mini-batching and shuffling during training
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=64, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

class DeepMLP(nn.Module):
    """Fully-connected classifier: 64 -> 128 -> 64 -> 10 with BatchNorm + Dropout."""
    def __init__(self, in_features=64, num_classes=10):
        super(DeepMLP, self).__init__()
        self.net = nn.Sequential(
            # Hidden block 1
            nn.Linear(in_features, 128),
            nn.BatchNorm1d(128), # stabilizes activations, speeds up training
            nn.ReLU(), # non-linear activation
            nn.Dropout(0.25), # regularization: randomly drop 25% of neurons

            # Hidden block 2
            nn.Linear(128, 64),
            nn.BatchNorm1d(64),
            nn.ReLU(),
            nn.Dropout(0.20),

            # Output layer: 10 raw logits (one per digit class)
            nn.Linear(64, num_classes)
        )

    def forward(self, x):
        return self.net(x)

model = DeepMLP().to(device) # move all parameters to CPU/GPU
print(model)
```

Listing 57: Convert scaled NumPy arrays into PyTorch tensors + DataLoaders

```
# ---- Training configuration ----
criterion = nn.CrossEntropyLoss() # softmax + NLL loss
optimizer = optim.Adam(model.parameters(), lr=0.003, weight_decay=1e-4) # weight decay = L2
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.5, patience=5)

epochs = 50
history = {'train_loss': [], 'val_loss': [], 'train_acc': [], 'val_acc': []}

for epoch in range(epochs):
    # ----- Training phase -----
    model.train() # enable Dropout/BatchNorm training behavior
    running_loss, correct, total = 0.0, 0, 0
    for inputs, labels in train_loader:
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad() # reset gradients from the previous step
        outputs = model(inputs) # forward pass
        loss = criterion(outputs, labels) # compute cross-entropy
```

```

        loss.backward()                # backpropagation (compute gradients)
        optimizer.step()               # update weights

        # Accumulate metrics (loss is averaged over the batch, so multiply by batch size)
        running_loss += loss.item() * inputs.size(0)
        _, preds = torch.max(outputs, 1) # predicted class = highest logit
        correct += (preds == labels).sum().item()
        total += labels.size(0)

    epoch_train_loss = running_loss / total
    epoch_train_acc = correct / total

    # ----- Validation phase -----
    model.eval()                      # disable Dropout, use running BatchNorm stats
    val_loss, val_correct, val_total = 0.0, 0, 0
    with torch.no_grad():              # no gradients needed for evaluation
        for inputs, labels in val_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            loss = criterion(outputs, labels)

            val_loss += loss.item() * inputs.size(0)
            _, preds = torch.max(outputs, 1)
            val_correct += (preds == labels).sum().item()
            val_total += labels.size(0)

    epoch_val_loss = val_loss / val_total
    epoch_val_acc = val_correct / val_total
    scheduler.step(epoch_val_loss)     # reduce LR if validation loss plateaus

    # Record metrics for plotting
    history['train_loss'].append(epoch_train_loss)
    history['val_loss'].append(epoch_val_loss)
    history['train_acc'].append(epoch_train_acc)
    history['val_acc'].append(epoch_val_acc)

    # Print a progress line every 10 epochs (and the first one)
    if (epoch + 1) % 10 == 0 or epoch == 0:
        print(f"Epoch [{epoch+1:02d}/{epochs}] "
              f"Train Loss: {epoch_train_loss:.4f} | Train Acc: {epoch_train_acc:.4f} || "
              f"Val Loss: {epoch_val_loss:.4f} | Val Acc: {epoch_val_acc:.4f}")

```

Listing 58: Training configuration

```

# ---- Final evaluation on the held-out test set ----
model.eval()
test_preds = []
test_targets = []

with torch.no_grad():
    for inputs, labels in test_loader:
        inputs = inputs.to(device)
        outputs = model(inputs)
        _, preds = torch.max(outputs, 1) # predicted digit
        test_preds.extend(preds.cpu().numpy()) # move back to CPU for metrics
        test_targets.extend(labels.numpy())

acc_pytorch = accuracy_score(test_targets, test_preds)
print(f"PyTorch Deep MLP Test Accuracy: {acc_pytorch:.4f}\n")
print("Detailed Classification Report:")
# Per-class precision/recall/F1 for all 10 digits
print(classification_report(test_targets, test_preds, digits=4))

```

Listing 59: Final evaluation on the held-out test set

## **Part VI**

# **Sequence Modeling: Recurrent Neural Networks**

## Experiment 14: Recurrent Neural Network (RNN)

**Dataset:** Synthetic 20-step sequences (1,200; label = positive sum) and Monthly Airline Passengers (12-month windows)

**Key parameters:** Vanilla RNN vs LSTM; 2 layers, hidden 64 (forecasting) / 32 (classification); Adam; MSE / BCE loss

**Headline result:** Forecasting: LSTM RMSE 44.12 vs RNN 69.93 passengers; Classification: LSTM 0.9450 vs RNN 0.9300

### Objective

Build recurrent networks for sequential data — sequence classification and time-series forecasting — and compare a vanilla RNN with an LSTM.

### Suggested Dataset

(a) Synthetic sequence benchmark: 1,200 sequences of length 20, label = 1 if the sum is positive (1,000 train / 200 test). (b) Monthly Airline Passengers (`data/airline_passengers.csv`): 144 months, 12-month sliding windows, chronological 80/20 split.

### Required Libraries

PyTorch, numpy, pandas, matplotlib (the manual's TensorFlow/Keras example is implemented equivalently in PyTorch, the framework installed in this environment).

### Brief Theory

A recurrent network processes a sequence one step at a time while maintaining a hidden state,

$$h_t = \tanh(W_{ih}x_t + b_{ih} + W_{hh}h_{t-1} + b_{hh}), \quad \hat{y} = W_{ho}h_T + b_o.$$

Training uses backpropagation through time (BPTT). Repeated multiplication by  $W_{hh}$  makes gradients vanish (or explode) over long sequences. The **LSTM** solves this with gates — forget  $f_t$ , input  $i_t$  and output  $o_t$  — and an additive cell state  $c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$ , which lets gradients flow across many steps.

### Procedure

1. Generate 1,200 labeled 20-step sequences (label = 1 if the sum is positive) and split 1,000/200.
2. Build vanilla RNN and LSTM classifiers (hidden 32, linear sigmoid output); train 10 epochs with Adam and binary cross-entropy; report test accuracy.
3. For forecasting, normalize the passenger series and build 12-month sliding windows with a chronological 80/20 split.
4. Train RNN and LSTM forecasters (2 layers, hidden 64, MSE, Adam, 120 epochs).
5. Report RMSE/MAE and plot the forecast trajectories.

### Reference Implementation

```
self.rnn = nn.RNN(1, hidden, batch_first=True) if cell == 'rnn' else nn.LSTM(1, hidden, batch_first=True)
self.fc = nn.Linear(hidden, 1) # classification output

self.lstm = nn.LSTM(input_size=1, hidden_size=64, num_layers=2, batch_first=True)
self.fc = nn.Linear(64, 1) # forecasting output
```


**Listing 60:** Recurrent model definitions (classification and forecasting)

**Table 18:** Sequence classification on the synthetic benchmark (200 test sequences).

| Model       | Test accuracy |
|-------------|---------------|
| Vanilla RNN | 0.9300        |
| LSTM        | <b>0.9450</b> |

**Table 19:** Airline passenger forecasting on the chronological test horizon (27 windows).

| Model       | RMSE (passengers) | MAE          |
|-------------|-------------------|--------------|
| Vanilla RNN | 69.93             | 62.36        |
| LSTM        | <b>44.12</b>      | <b>36.78</b> |

**Figure 8:** Airline passenger forecasting: ground truth versus the RNN and LSTM forecasts on the test horizon.



## Evaluation / Expected Output

The LSTM outperforms the vanilla RNN on both tasks: +1.5 accuracy points on classification and about 37% lower RMSE on forecasting. The gap is larger on forecasting because that task needs 12-step seasonal memory, exactly where gating prevents vanishing gradients. On the easy synthetic benchmark both models do well because the decision rule ( $\text{sum} > 0$ ) is nearly linear.

## Viva / Discussion Questions

**Why does an RNN have memory?** Its hidden state is carried across time steps, so each output depends on previous inputs, not only the current one.

**What is a hidden state?** The internal vector summarizing the sequence seen so far; it is updated at every step and used to produce outputs.

**Why are LSTM/GRU often preferred for long dependencies?** Their gates control what is stored or forgotten and the additive cell state lets gradients flow across many steps, avoiding the exponential decay of vanilla RNN gradients.

**Why is the train/test split chronological?** Shuffling sequences would leak future values into training and produce an optimistic, invalid evaluation.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Scaling + regression metrics ----
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import root_mean_squared_error, mean_absolute_error, accuracy_score

# ---- PyTorch for the recurrent models ----
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

# Aesthetics setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 5)
plt.rcParams['figure.dpi'] = 110

# Reproducibility
torch.manual_seed(42)
torch.cuda.manual_seed_all(42) # reproducible GPU training
np.random.seed(42)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"PyTorch version: {torch.__version__} on {device}")

# ---- Load the monthly airline passengers time series ----
df_air = pd.read_csv('../data/airline_passengers.csv')
print("Dataset Head:")
display(df_air.head())

# Extract the target series as float32 (shape: N x 1)
passengers = df_air['Passengers'].values.astype(np.float32).reshape(-1, 1)
# ---- Normalize to [0, 1] (RNNs train much better on small inputs) ----
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(passengers)

# ---- Function to build supervised sliding-window sequences ----
# For each time step t: X = values [t-seq_len ... t-1], y = value at t
def create_sequences(data, seq_length=12):
    xs, ys = [], []
    for i in range(len(data) - seq_length):
        x = data[i:(i + seq_length)]
        y = data[i + seq_length]
        xs.append(x)
```

```

        ys.append(y)
    return np.array(xs), np.array(ys)

SEQ_LENGTH = 12 # 12-month lookback = one full annual seasonal cycle
X_seq, y_seq = create_sequences(scaled_data, seq_length=SEQ_LENGTH)

# ---- Temporal split: 80% train, 20% test WITHOUT shuffling (time order matters) ----
train_size = int(len(X_seq) * 0.80)
X_train, X_test = X_seq[:train_size], X_seq[train_size:]
y_train, y_test = y_seq[:train_size], y_seq[train_size:]

print(f"Total sequences created: {len(X_seq)}")
print(f"Training sequences:      {X_train.shape[0]} (Shape: {X_train.shape})")
print(f"Testing sequences:       {X_test.shape[0]} (Shape: {X_test.shape})")

# ---- DataLoader for mini-batch training (shuffle only the training set) ----
train_loader = DataLoader(
    TensorDataset(torch.tensor(X_train, dtype=torch.float32), torch.tensor(y_train, dtype=torch.float32)),
    batch_size=16,
    shuffle=True
)

```

Listing 61: Core numerical and plotting libraries

```

class VanillaRNN(nn.Module):
    """Elman RNN: 2 stacked recurrent layers + a linear output layer."""
    def __init__(self, input_size=1, hidden_size=64, num_layers=2):
        super(VanillaRNN, self).__init__()
        self.rnn = nn.RNN(input_size, hidden_size, num_layers, batch_first=True)
        self.fc = nn.Linear(hidden_size, 1) # map final hidden state to a scalar forecast

    def forward(self, x):
        out, _ = self.rnn(x) # out: (batch, seq_len, hidden)
        out = self.fc(out[:, -1, :]) # use ONLY the last time step's hidden state
        return out

class LSTMModel(nn.Module):
    """LSTM: same interface, but gated cells preserve long-term information."""
    def __init__(self, input_size=1, hidden_size=64, num_layers=2):
        super(LSTMModel, self).__init__()
        self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True)
        self.fc = nn.Linear(hidden_size, 1)

    def forward(self, x):
        out, _ = self.lstm(x)
        out = self.fc(out[:, -1, :])
        return out

# Instantiate both models on the selected device so they can be compared fairly
rnn_model = VanillaRNN().to(device)
lstm_model = LSTMModel().to(device)
print("Models instantiated successfully!")

```

Listing 62: Vanilla RNN and LSTM model definitions

```

# ---- Generic training function shared by both architectures ----
def train_model(model, loader, epochs=120, lr=0.005):
    criterion = nn.MSELoss() # regression -> mean squared error
    optimizer = optim.Adam(model.parameters(), lr=lr) # Adam optimizer
    losses = []

    for epoch in range(epochs):
        model.train()
        total_loss = 0.0
        for x_b, y_b in loader:
            x_b, y_b = x_b.to(device), y_b.to(device)

            optimizer.zero_grad() # clear previous gradients
            pred = model(x_b) # forward pass through the recurrent net
            loss = criterion(pred, y_b) # MSE between forecast and truth
            loss.backward() # backpropagation through time (BPTT)
            optimizer.step() # update weights

        total_loss += loss.item() * len(x_b) # de-average the batch loss

```

```

        epoch_loss = total_loss / len(loader.dataset)
        losses.append(epoch_loss)          # track loss per epoch for plotting

    return losses

# ---- Train both architectures with identical hyperparameters ----
print("Training Vanilla RNN...")
rnn_losses = train_model(rnn_model, train_loader, epochs=120, lr=0.005)

print("Training LSTM Model...")
lstm_losses = train_model(lstm_model, train_loader, epochs=120, lr=0.005)
print("Training completed!")

```

Listing 63: Generic training function shared by both architectures

```

# ---- Evaluate both models on the test horizon ----
rnn_model.eval()
lstm_model.eval()

with torch.no_grad():    # inference only - no gradients
    x_test_t = torch.tensor(X_test, dtype=torch.float32).to(device)
    pred_rnn_scaled = rnn_model(x_test_t).cpu().numpy()    # normalized predictions
    pred_lstm_scaled = lstm_model(x_test_t).cpu().numpy()

# ---- Invert scaling to get real passenger counts ----
y_test_true = scaler.inverse_transform(y_test)
pred_rnn = scaler.inverse_transform(pred_rnn_scaled)
pred_lstm = scaler.inverse_transform(pred_lstm_scaled)

# ---- Compute error metrics in the original unit (passengers) ----
rmse_rnn = root_mean_squared_error(y_test_true, pred_rnn)
mae_rnn = mean_absolute_error(y_test_true, pred_rnn)

rmse_lstm = root_mean_squared_error(y_test_true, pred_lstm)
mae_lstm = mean_absolute_error(y_test_true, pred_lstm)

# ---- Summary table (lower errors are better) ----
eval_df = pd.DataFrame([
    {"Model": "Vanilla RNN", "RMSE (Passengers)": rmse_rnn, "MAE (Passengers)": mae_rnn},
    {"Model": "LSTM Network", "RMSE (Passengers)": rmse_lstm, "MAE (Passengers)": mae_lstm}
]).set_index("Model")

print("Time Series Forecasting Performance Comparison:")
display(eval_df.style.highlight_min(subset=['RMSE (Passengers)', 'MAE (Passengers)'], color='lightgreen'))

# ---- Plot ground truth vs. both forecast trajectories ----
# Map each test target back to its true position on the month axis:
# sequence i predicts month (i + SEQ_LENGTH)
test_indices = range(train_size + SEQ_LENGTH, len(passengers))

plt.figure(figsize=(12, 6))
# Full historical series in the background
plt.plot(range(len(passengers)), passengers, label='Ground Truth Historical', color='black', alpha=0.5, lw=1.5)

# Test-period ground truth
plt.plot(test_indices, y_test_true, label='Ground Truth (Test Horizon)', color='black', lw=2.5)

# Model forecasts
plt.plot(test_indices, pred_rnn, label=f'Vanilla RNN (RMSE: {rmse_rnn:.1f})', color='crimson', lw=2, linestyle='--',
        ⇐ ')
plt.plot(test_indices, pred_lstm, label=f'LSTM Network (RMSE: {rmse_lstm:.1f})', color='teal', lw=2.2)

# Vertical line marking where the train period ends and test begins
plt.axvline(x=train_size + SEQ_LENGTH, color='gray', linestyle=':', label='Train / Test Cutoff')
plt.title("Airline Passengers Sequence Forecasting: Vanilla RNN vs LSTM", fontsize=14, fontweight='bold')
plt.xlabel("Month Index", fontsize=12)
plt.ylabel("Number of Passengers", fontsize=12)
plt.legend()
plt.tight_layout()
plt.show()

```

Listing 64: Evaluate both models on the test horizon

```
# ---- Sequence classification benchmark (manual experiment) ----
rng = np.random.default_rng(42)
X_cls = rng.normal(size=(1200, 20, 1)).astype("float32")
y_cls = (X_cls.sum(axis=1).ravel() > 0).astype("float32")

X_tr_c, X_te_c = X_cls[:1000], X_cls[1000:]
y_tr_c, y_te_c = y_cls[:1000], y_cls[1000:]
print(f"Sequence classification data: train {X_tr_c.shape}, test {X_te_c.shape}")

class RNNClassifier(nn.Module):
    """Vanilla RNN or LSTM classifier over a 20-step sequence."""
    def __init__(self, cell='rnn', hidden=32):
        super().__init__()
        self.rnn = nn.RNN(1, hidden, batch_first=True) if cell == 'rnn' else nn.LSTM(1, hidden, batch_first=True)
        self.fc = nn.Linear(hidden, 1)

    def forward(self, x):
        out, _ = self.rnn(x)
        return self.fc(out[:, -1, :]).squeeze(-1)

def train_classifier(model, epochs=10):
    """Train a sequence classifier and return test accuracy."""
    model = model.to(device)
    optimizer = optim.Adam(model.parameters(), lr=0.005)
    loss_fn = nn.BCEWithLogitsLoss()
    loader = DataLoader(TensorDataset(torch.tensor(X_tr_c), torch.tensor(y_tr_c)), batch_size=32, shuffle=True)
    for epoch in range(epochs):
        model.train()
        for xb, yb in loader:
            xb, yb = xb.to(device), yb.to(device)
            optimizer.zero_grad()
            loss = loss_fn(model(xb), yb)
            loss.backward()
            optimizer.step()
        model.eval()
        with torch.no_grad():
            logits = model(torch.tensor(X_te_c).to(device))
            preds = (torch.sigmoid(logits) > 0.5).cpu().numpy().astype(int)
        return accuracy_score(y_te_c.astype(int), preds)

acc_rnn_cls = train_classifier(RNNClassifier('rnn'))
acc_lstm_cls = train_classifier(RNNClassifier('lstm'))
print(f"Sequence classification test accuracy: Vanilla RNN {acc_rnn_cls:.4f} | LSTM {acc_lstm_cls:.4f}")
```

Listing 65: Sequence classification benchmark (manual experiment)

## **Part VII**

# **Unsupervised Neural Learning: Self-Organizing Maps**

## Experiment 15: Self-Organizing Map (SOM)

**Dataset:** UCI Wine: 178 samples, 13 chemical features, 3 cultivars (labels used only for visualization)

**Key parameters:**  $12 \times 12$  grid, Gaussian neighborhood,  $\sigma = 1.5$ , learning rate 0.5, PCA initialization, 5,000 iterations

**Headline result:** Quantization error  $0.4447 \rightarrow 0.1965$ ; topographic error 0.0393

### Objective

Use a Self-Organizing Map to map high-dimensional observations onto a low-dimensional grid while preserving neighborhood structure.

### Suggested Dataset

UCI Wine: 178 samples, 13 chemical features, 3 cultivars (labels used only for visualization). Features are MinMax-scaled to  $[0, 1]$ .

### Required Libraries

minisom, numpy, pandas, matplotlib.

### Brief Theory

A Self-Organizing Map is a grid of neurons, each with a weight vector  $w_j$  in the input space. For each training sample the algorithm performs three steps:

1. **competition:** find the Best Matching Unit  $c = \arg \min_j \|x - w_j\|$ ;
2. **cooperation:** compute the Gaussian neighborhood  $h_{cj}(t) = \exp(-\|r_c - r_j\|^2 / (2\sigma(t)^2))$  on grid coordinates  $r$ ;
3. **adaptation:** move the weights toward the input,  $w_j \leftarrow w_j + \alpha(t) h_{cj}(t) (x - w_j)$ .

The learning rate  $\alpha(t)$  and radius  $\sigma(t)$  decay over time: early large neighborhoods organize the global layout, later small updates refine details. The result is a topology-preserving 2D projection. Quality is measured by the **quantization error** (average distance from samples to their BMU) and the **topographic error** (fraction of samples whose two best neurons are not neighbors). The **U-matrix** visualizes cluster boundaries as ridges of high inter-neuron distance.

### Procedure

1. Load the wine data and MinMax-scale to  $[0, 1]$ .
2. Create a  $12 \times 12$  SOM with a Gaussian neighborhood ( $\sigma = 1.5$ , learning rate 0.5) and initialize the weights with PCA.
3. Train for 5,000 random samples.
4. Compute the quantization and topographic errors.
5. Plot the U-matrix with each sample projected onto its BMU (marker = true cultivar).
6. Discuss grid size and  $\sigma$ .

### Reference Implementation


```
som = MiniSom(x=12, y=12, input_len=13, sigma=1.5, learning_rate=0.5,
              neighborhood_function='gaussian', random_seed=42)
som.pca_weights_init(X)
som.train_random(data=X, num_iteration=5000)
qe = som.quantization_error(X); te = som.topographic_error(X)
```

**Listing 66:** Creating, initializing and training the Kohonen map

## Evaluation / Expected Output

**Table 20:** SOM quality metrics on the Wine data.

| Metric                     | Value                                   |
|----------------------------|-----------------------------------------|
| Initial quantization error | 0.4447                                  |
| Final quantization error   | <b>0.1965</b>                           |
| Topographic error          | <b>0.0393</b> ( $\approx 4\%$ foldings) |



**Figure 9:** SOM U-matrix with the wine samples projected onto their BMUs; marker shape and color encode the true cultivar.

The map halved its quantization error during training and preserves topology well (topographic error below 0.05). On the U-matrix the three wine cultivars occupy contiguous regions separated by bright ridges, even though the SOM never saw the labels — evidence of unsupervised structure discovery. PCA initialization gave a well-unfolded map from the start.

## Viva / Discussion Questions

**SOM vs K-means?** Both compress data with prototypes, but K-means only finds centroids while SOM arranges them on a grid so neighboring neurons represent similar inputs (topology preservation).

**What is a BMU?** The neuron whose weight vector is closest to the input — the winner of the competition.

**Why is SOM called a topology-preserving map?** The neighborhood update ensures nearby grid positions end up with similar weight vectors, so high-dimensional neighborhoods are approximately preserved in 2D.

**Why initialize with PCA?** Random initialization can create folded maps; PCA starts the map spread along the data's principal directions, giving faster and more faithful convergence.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
```

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Dataset + scaling ----
from sklearn.datasets import load_wine
from sklearn.preprocessing import MinMaxScaler
# ---- MiniSom: lightweight Kohonen Self-Organizing Map implementation ----
from minisom import MiniSom

# Styling setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110

np.random.seed(42)
print("MiniSom and libraries loaded successfully!")

# ---- Load the Wine recognition dataset ----
wine = load_wine()
X_raw = wine.data          # 178 samples x 13 chemical features
y = wine.target            # 3 wine cultivars (only used for visualization - SOM is unsupervised)
feature_names = wine.feature_names
target_names = wine.target_names

print(f"Samples: {X_raw.shape[0]}, Features: {X_raw.shape[1]}, Classes: {target_names.tolist()}")

# ---- MinMax scaling to [0, 1] ----
# SOM weight vectors live in the same range as the inputs, so scaling keeps learning stable.
scaler = MinMaxScaler()
X = scaler.fit_transform(X_raw)

```

Listing 67: Core numerical and plotting libraries

```

# ---- Configure the Kohonen grid ----
grid_x, grid_y = 12, 12 # 12x12 grid = 144 neurons (fewer than the 178 samples)
input_len = X.shape[1]  # 13 chemical features per neuron weight vector

som = MiniSom(
    x=grid_x,
    y=grid_y,
    input_len=input_len,
    sigma=1.5,                # initial neighborhood radius
    learning_rate=0.5,        # initial learning rate
    neighborhood_function='gaussian', # smooth Gaussian neighborhood kernel
    random_seed=42
)

# ---- Initialize weights with PCA for topologically faithful, faster convergence ----
som.pca_weights_init(X)
print("Initial Quantization Error:", som.quantization_error(X))

# ---- Train: 5000 random samples presented to the map ----
som.train_random(data=X, num_iteration=5000, verbose=False)
final_qe = som.quantization_error(X) # average distance from samples to their BMU
final_te = som.topographic_error(X)  # fraction of samples whose 2nd BMU is not a grid neighbor

print(f"Training complete after 5,000 iterations!")
print(f" - Final Quantization Error (mean distance to BMU): {final_qe:.4f}")
print(f" - Final Topographic Error (proportion of foldings): {final_te:.4f}")

```

Listing 68: Configure the Kohonen grid

```

# ---- Plot the U-Matrix (unified distance map) ----
# Each cell = average distance between a neuron and its immediate grid neighbors.
# Dark valleys = dense clusters; bright ridges = boundaries between clusters.
plt.figure(figsize=(10, 8))
plt.pcolor(som.distance_map().T, cmap='bone_r', alpha=0.9)
cbar = plt.colorbar()
cbar.set_label("Normalized Inter-Neuron Distance (U-Matrix)", fontsize=11)

# ---- Project every sample onto its Best Matching Unit (BMU) ----

```



```
markers = ['o', 's', '^']
colors = ['red', 'green', 'blue']

for idx, x_vec in enumerate(X):
    w = som.winner(x_vec)  # grid coordinates of the best matching neuron for this sample
    plt.plot(
        w[0] + 0.5,        # +0.5 centers the marker inside the pcolor cell
        w[1] + 0.5,
        markers[y[idx]],   # marker shape encodes the true cultivar
        markerfacecolor='None',
        markeredgecolor=colors[y[idx]],
        markersize=10,
        markeredgewidth=2
    )

# Build a legend for the three cultivars
for c_idx, c_name in enumerate(target_names):
    plt.plot([], [], marker=markers[c_idx], color=colors[c_idx], linestyle='None',
             label=f'Cultivar {c_name}', markersize=9, markeredgewidth=2)

plt.title("SOM U-Matrix with Projected Wine Cultivars", fontsize=14, fontweight='bold')
plt.legend(bbox_to_anchor=(1.25, 1), loc='upper left')
plt.tight_layout()
plt.show()
```

**Listing 69:** Plot the U-Matrix (unified distance map)

## **Part VIII**

# **Probabilistic Sequence Modeling: Hidden Markov Models**

## Experiment 16: Hidden Markov Model (HMM)

**Dataset:** DAX stock index 1991–1998 (`data/stock_index.csv`, `real EuStockMarkets`); 1,859 trading days  
**Key parameters:** 3 Gaussian states, full covariance, 200 Baum–Welch (EM) iterations; Viterbi decoding  
**Headline result:** EM converged; log-likelihood 6019.79; volatility-sorted Bull / Sideways / Bear regimes

### Objective

Model a sequence using hidden states, transition probabilities and emission distributions; decode the most likely hidden-state path with the Viterbi algorithm.

### Suggested Dataset

Real DAX stock index daily closes 1991–1998 (`data/stock_index.csv`, `R EuStockMarkets`), 1,859 trading days. The observation sequence is the daily return; the close price is used only for plotting.

### Required Libraries

hmmlearn, numpy, pandas, matplotlib.

### Brief Theory

A Hidden Markov Model is a doubly stochastic process:

- a hidden state sequence  $z_t$  following a first-order Markov chain,  $P(z_t \mid z_{t-1}, \dots, z_1) = P(z_t \mid z_{t-1})$ , summarized by the transition matrix  $A_{ij} = P(z_{t+1} = s_j \mid z_t = s_i)$ ;
- observations emitted from the active state,  $x_t \mid (z_t = s_k) \sim \mathcal{N}(\mu_k, \sigma_k^2)$  for a Gaussian HMM.

Three classical problems are solved here: **evaluation** (the forward algorithm gives  $P(X \mid \lambda)$ ), **decoding** (the Viterbi algorithm finds the most likely state path by dynamic programming) and **learning** (Baum–Welch EM alternates a forward–backward E-step with parameter re-estimation until the log-likelihood converges).

### Procedure

1. Load the DAX series and build the daily-return observation vector.
2. Fit a 3-state Gaussian HMM with full covariance and 200 EM iterations.
3. Check convergence and the model log-likelihood.
4. Decode the state path with Viterbi.
5. Sort the states by volatility so they map to Bull/Sideways/Bear; reorder the transition matrix accordingly.
6. Plot the transition heatmap and the price series colored by regime.

### Reference Implementation

```
hmm_model = GaussianHMM(n_components=3, covariance_type="full", n_iter=200, random_state=42)
hmm_model.fit(returns)
hidden_states = hmm_model.predict(returns) # Viterbi
state_order = np.argsort([np.sqrt(hmm_model.covars_[i][0][0]) for i in range(3)])
```



**Listing 70:** Fitting, decoding and volatility-sorting the HMM

### Evaluation / Expected Output

The model extracts interpretable, persistent market regimes from real data without labels. The colored DAX chart tracks real episodes: the 1990s bull phases, corrections such as 1994, and the high-volatility cluster around the 1998 crisis. Its

**Table 21:** Gaussian HMM on the DAX index (1,859 trading days, 1991–1998).

| Item                 | Value                                              |
|----------------------|----------------------------------------------------|
| EM converged         | True                                               |
| Model log-likelihood | <b>6019.79</b>                                     |
| Regimes              | Bull (low vol), Sideways (medium), Bear (high vol) |

**Figure 10:** Estimated transition-probability matrix after sorting the states by volatility; the dominant diagonal indicates regime persistence.**Figure 11:** DAX index 1991–1998 colored by the Viterbi-decoded regime (top) and the decoded state sequence over time (bottom).

main assumptions are the Markov property and Gaussian emissions; return fat tails and duration dependence motivate extensions such as  $t$ -emissions or semi-Markov models.

## Viva / Discussion Questions

**What is hidden in an HMM?** The state sequence: we observe emissions (returns) but not the underlying regime that generated them.

**Viterbi vs forward algorithm?** Viterbi returns the single most likely state path (a maximum over paths); the forward algorithm computes the total probability  $P(X | \lambda)$  by summing over all paths.

**What is the Markov assumption?** The future state depends on the past only through the present state:  $P(z_t | z_{1:t-1}) = P(z_t | z_{t-1})$ .

**Why sort the states by volatility?** So that regime labels (Bull/Sideways/Bear) are stable and interpretable across runs.

**Why model returns and not prices?** Returns are approximately stationary; prices trend and violate the Gaussian emission assumption.

## Complete Code Listing

```
# ---- Core libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from hmmlearn.hmm import GaussianHMM

# Styling setup
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110
np.random.seed(42)

# ---- Load the real DAX stock index (daily closes 1991-1998) ----
df_market = pd.read_csv('../data/stock_index.csv')
dates = pd.to_datetime(df_market['Date'])
print(f"Trading days: {len(df_market)} | {dates.iloc[0].date()} to {dates.iloc[-1].date()}")
print(df_market.head(3).to_string(index=False))

returns = df_market['Daily_Return'].values.reshape(-1, 1) # HMM observations
prices = df_market['Close'].values # only used for visualization
```

**Listing 71: Core libraries**

```
# ---- Fit a 3-state Gaussian HMM with the Baum-Welch (EM) algorithm ----
n_components = 3
hmm_model = GaussianHMM(
    n_components=n_components, # Bull / Sideways / Bear regimes
    covariance_type="full", # each state has its own full covariance Gaussian emission
    n_iter=200, # EM iterations
    random_state=42,
    verbose=False
)

hmm_model.fit(returns)
print("EM Training Converged:", hmm_model.monitor_.converged)
print(f"Model Log-Likelihood: {hmm_model.score(returns):.2f}")

# ---- Decode the most likely hidden state sequence with the Viterbi algorithm ----
hidden_states = hmm_model.predict(returns)

# ---- Sort regimes by volatility so labels have a consistent economic meaning ----
# Regime 0: Low Volatility (Bull) | Regime 1: Moderate (Sideways) | Regime 2: High (Bear/Crisis)
volatilities = [np.sqrt(hmm_model.covars_[i][0][0]) for i in range(n_components)]
state_order = np.argsort(volatilities) # ascending volatility
state_mapping = {old_state: new_state for new_state, old_state in enumerate(state_order)}

sorted_hidden_states = np.array([state_mapping[s] for s in hidden_states]) # relabel sequence
sorted_means = [hmm_model.means_[old][0] for old in state_order] # reordered means
```

```
sorted_stdvs = [np.sqrt(hmm_model.covars_[old][0][0]) for old in state_order] # reordered std devs

# ---- Reorder the transition matrix to match the sorted regimes ----
sorted_transmat = np.zeros((n_components, n_components))
for i in range(n_components):
    for j in range(n_components):
        sorted_transmat[i, j] = hmm_model.transmat_[state_order[i], state_order[j]]

regime_names = ["Bull (Low Vol)", "Sideways (Med Vol)", "Bear (High Vol)"]
```

**Listing 72:** Fit a 3-state Gaussian HMM with the Baum-Welch (EM) algorithm

```
# ---- Transition probability matrix ----
fig, ax = plt.subplots(figsize=(6, 4.5))
sns.heatmap(sorted_transmat, annot=True, fmt='.3f', cmap='Blues', ax=ax,
            xticklabels=regime_names, yticklabels=regime_names, cbar=False)
ax.set_title("HMM Transition Probability Matrix ($A_{ij}$)", fontweight='bold')
ax.set_xlabel("Transition To Regime")
ax.set_ylabel("Current Regime")
plt.tight_layout()
plt.show()

# ---- Viterbi-decoded regimes over the price series ----
colors = ['forestgreen', 'royalblue', 'crimson']
x = np.arange(len(prices))

fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(13, 7), sharex=True, gridspec_kw={'height_ratios': [3, 1]})

# Top: color every consecutive price segment by its decoded regime
for i in range(len(prices) - 1):
    ax1.plot([x[i], x[i + 1]], [prices[i], prices[i + 1]], color=colors[sorted_hidden_states[i]], lw=1.2)
for i, name in enumerate(regime_names):
    ax1.plot([], [], color=colors[i], label=name, lw=2.5)
ax1.set_title("DAX Index 1991-1998 Colored by Decoded HMM Regime (Viterbi Path)", fontsize=13, fontweight='bold')
ax1.set_ylabel("Index Level")
ax1.legend(loc='upper left')

# Bottom: raw sequence of decoded states over time
ax2.scatter(x, sorted_hidden_states, c=[colors[s] for s in sorted_hidden_states], s=8, marker='|')
ax2.set_yticks(range(n_components))
ax2.set_yticklabels(regime_names)
ax2.set_xlabel("Trading Day Index")

# Label the x-axis with years (real dates)
year_pos = np.where(dates.dt.year.diff().fillna(0).values != 0)[0]
ax2.set_xticks(year_pos)
ax2.set_xticklabels(dates.dt.year.iloc[year_pos].astype(str))

plt.tight_layout()
plt.show()
```

**Listing 73:** Transition probability matrix

## **Part IX**

# **Margin-Based Learning: Support Vector Machines**

## Experiment 17: Support Vector Machine (SVM)

**Dataset:** Breast Cancer ( $569 \times 30$ ; 2 features for the kernel view, all 30 for accuracy) and Old Faithful geyser (`data/geyser.csv`, 272 eruptions)

**Key parameters:** SVC: RBF kernel,  $C = 1.0$ ,  $\gamma = \text{'scale'}$ ; SVR: RBF,  $C = 10$ ,  $\varepsilon = 0.3$

**Headline result:** SVC accuracy 0.9790 (96/426 support vectors); SVR  $R^2$  0.8968 (103/272 support vectors)

### Objective

Train SVM classifiers and investigate the effects of kernel,  $C$  and  $\gamma$ ; demonstrate Support Vector Regression with an  $\varepsilon$ -insensitive tube.

### Suggested Dataset

(a) Breast Cancer ( $569 \times 30$ ): two features (mean radius, mean texture) for the 2D kernel view and all 30 features for the accuracy benchmark. (b) Old Faithful geyser (`data/geyser.csv`, 272 eruptions) for SVR.

### Required Libraries

scikit-learn, numpy, pandas, matplotlib.

### Brief Theory

SVC (soft margin) solves

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + C \sum_i \xi_i \quad \text{s.t.} \quad y_i (w^T \phi(x_i) + b) \geq 1 - \xi_i, \quad \xi_i \geq 0.$$

The dual problem depends only on inner products  $K(x_i, x_j)$ , which enables the **kernel trick**: compute similarities in a high-dimensional feature space without mapping the data there explicitly. The solution is sparse — only points with  $\alpha_i > 0$  (**support vectors**) define the boundary. Kernels used here: linear  $x^T z$ , polynomial  $(\gamma x^T z + r)^d$  and RBF  $\exp(-\gamma \|x - z\|^2)$ .  $C$  penalizes margin violations;  $\gamma$  controls the RBF width. **SVR** applies the same machinery to regression with an  $\varepsilon$ -insensitive tube: residuals within  $\pm \varepsilon$  cost nothing and only points outside the tube become support vectors.

### Procedure

1. Scale the features.
2. Train SVC with linear, polynomial and RBF kernels on two breast-cancer features; plot the boundaries, margins and support vectors.
3. Train the RBF SVC on all 30 features; report accuracy and the support-vector count.
4. Fit SVR (RBF,  $C = 10$ ,  $\varepsilon = 0.3$ ) to the Old Faithful data (waiting time  $\rightarrow$  eruption duration); plot the fitted curve and the tube.
5. Report  $R^2$ , RMSE and the support-vector share.

### Reference Implementation



```
clf = SVC(kernel=k_name, C=1.0, gamma='scale', degree=3, random_state=42).fit(X_2d_scaled, y_2d)
svc_rbf = SVC(kernel='rbf', C=1.0, random_state=42).fit(X_train_c_s, y_train_c)
svr_model = SVR(kernel='rbf', C=10.0, epsilon=0.3, gamma='scale').fit(X_svr_scaled, y_svr)
```

**Listing 74:** SVC per kernel and SVR with an epsilon-insensitive tube



**Table 22:** SVM results on the two real datasets.

| Experiment             | Metric          | Value              |
|------------------------|-----------------|--------------------|
| SVC (RBF, 30 features) | Test accuracy   | <b>0.9790</b>      |
| SVC (RBF, 30 features) | Support vectors | 96 of 426 (22.5%)  |
| SVR (Old Faithful)     | $R^2$           | 0.8968             |
| SVR (Old Faithful)     | RMSE            | 0.366 minutes      |
| SVR (Old Faithful)     | Support vectors | 103 of 272 (37.9%) |

**Figure 12:** SVC decision boundaries, margins and support vectors on two real breast-cancer features for the linear, polynomial and RBF kernels.**Figure 13:** Support Vector Regression on the Old Faithful geyser with the  $\epsilon = 0.3$  tube; highlighted points are support vectors.

## Evaluation / Expected Output

The 2D kernel figure shows the linear kernel underfitting the real diagnostic data while the polynomial and RBF kernels produce curved boundaries. The RBF SVC reaches 0.9790 accuracy with only 96 of 426 training points as support vectors — about 77% of the data could be discarded without changing the classifier. The geyser scatter is bimodal (short and long eruptions); the RBF SVR smoothly bridges the two modes and the tube leaves most residuals unpenalized.

## Viva / Discussion Questions

**What is a support vector?** A training point on or inside the margin ( $\alpha > 0$ ) that alone defines the decision boundary.

**Effect of  $C$ ?** Large  $C$  penalizes violations strongly (narrow margin, possible overfitting); small  $C$  allows a wider margin and more violations.

**Effect of  $\gamma$ ?** Large  $\gamma$  makes each point's influence local (wiggly boundary); small  $\gamma$  makes the RBF smoother, approaching a linear model.

**What is the kernel trick?** Computing inner products in a high-dimensional feature space without explicitly mapping the points there.

**What is the  $\varepsilon$ -tube in SVR?** The zero-loss zone around the fitted function; only points outside it become support vectors. A wider tube (0.3 vs. 0.1) reduced the support-vector count with the same  $R^2$ .

**Why does SVM require careful scaling?** Margins and kernels use Euclidean distances; unscaled features with larger ranges dominate them and degrade the fit.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- Dataset + model selection ----
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
# ---- SVM models: classification and regression ----
from sklearn.svm import SVC, SVR
# ---- Evaluation metrics ----
from sklearn.metrics import r2_score, root_mean_squared_error

# Styling setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.dpi'] = 110

np.random.seed(42)
print("SVM modules loaded successfully!")
```

**Listing 75:** Core numerical and plotting libraries

```
# ---- 2D real data for visualizing margins and kernels ----
# Two diagnostic measurements from the Breast Cancer dataset (malignant vs benign).
cancer_2d = load_breast_cancer()
X_2d = cancer_2d.data[:, [0, 1]] # mean radius and mean texture
y_2d = cancer_2d.target
feature_names_2d = ['mean radius', 'mean texture']

scaler_2d = StandardScaler()
X_2d_scaled = scaler_2d.fit_transform(X_2d)

kernels = ['linear', 'poly', 'rbf']
fig, axes = plt.subplots(1, 3, figsize=(18, 5))

# Coordinate grid used to draw the decision boundary as a contour
x_min, x_max = X_2d_scaled[:, 0].min() - 0.5, X_2d_scaled[:, 0].max() + 0.5
y_min, y_max = X_2d_scaled[:, 1].min() - 0.5, X_2d_scaled[:, 1].max() + 0.5
```

```

xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200), np.linspace(y_min, y_max, 200))

for ax, k_name in zip(axes, kernels):
    # Train one SVC per kernel on the same real data
    clf = SVC(kernel=k_name, C=1.0, gamma='scale', degree=3, random_state=42)
    clf.fit(X_2d_scaled, y_2d)

    # Signed distance to the hyperplane for every grid point
    Z = clf.decision_function(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)

    # Filled contour = decision regions; solid line = boundary (Z=0); dashed = margins (Z=±1)
    ax.contourf(xx, yy, Z, levels=np.linspace(Z.min(), Z.max(), 20), cmap='coolwarm', alpha=0.3)
    ax.contour(xx, yy, Z, levels=[-1.0, 0.0, 1.0], linestyles=['--', '-', '--'], colors=['k', 'k', 'k'],
    ↪ linewidths=[1.5, 2.5, 1.5])

    # Scatter the data points (colored by class)
    ax.scatter(X_2d_scaled[:, 0], X_2d_scaled[:, 1], c=y_2d, cmap='coolwarm', edgecolors='k', s=25, alpha=0.8)

    # Highlight the support vectors (the points that define the margin)
    sv = clf.support_vectors_
    ax.scatter(sv[:, 0], sv[:, 1], s=90, facecolors='none', edgecolors='gold', linewidths=2.0, label=f'SV ({len(sv)
    ↪}))')

    ax.set_title(f"SVC with {k_name.capitalize()} Kernel\n(Acc: {clf.score(X_2d_scaled, y_2d):.3f}, SVs: {len(sv)
    ↪})", fontweight='bold')
    ax.set_xlabel(feature_names_2d[0])
    ax.set_ylabel(feature_names_2d[1])
    ax.legend(loc='lower right')

plt.suptitle("SVC Decision Boundaries on Real Breast Cancer Data (2 features)", fontsize=15, fontweight='bold')
plt.tight_layout()
plt.show()

```

Listing 76: 2D real data for visualizing margins and kernels

```

# ---- SVC on a real dataset: Breast Cancer (RBF kernel) ----
cancer = load_breast_cancer()
X_train_c, X_test_c, y_train_c, y_test_c = train_test_split(
    cancer.data, cancer.target, test_size=0.25, random_state=42, stratify=cancer.target
)
scaler_c = StandardScaler()
X_train_c_s = scaler_c.fit_transform(X_train_c)
X_test_c_s = scaler_c.transform(X_test_c)

svc_rbf = SVC(kernel='rbf', C=1.0, random_state=42)
svc_rbf.fit(X_train_c_s, y_train_c)
print(f"SVC (RBF) test accuracy on Breast Cancer: {svc_rbf.score(X_test_c_s, y_test_c):.4f}")
print(f"Support vectors used: {len(svc_rbf.support_)} of {len(X_train_c)} training samples")

```

Listing 77: SVC on a real dataset: Breast Cancer (RBF kernel)

```

# ---- Real 1D regression: Old Faithful geyser ----
# x = waiting time before an eruption (minutes), y = eruption duration (minutes)
geyser = pd.read_csv('../data/geyser.csv')
X_svr = geyser['waiting'].values.reshape(-1, 1)
y_svr = geyser['eruptions'].values
print(f"Old Faithful: {len(geyser)} eruptions | waiting {X_svr.min():.1f}-{X_svr.max():.1f} min | duration {y_svr.min():.1f}
    ↪ {y_svr.max():.1f} min")

# Scale the feature (SVR is distance-based); the target stays in minutes for interpretability
scaler_svr = StandardScaler()
X_svr_scaled = scaler_svr.fit_transform(X_svr)

# epsilon defines the width of the "no penalty" tube around the prediction (in minutes)
epsilon_val = 0.3
svr_model = SVR(kernel='rbf', C=10.0, epsilon=epsilon_val, gamma='scale')
svr_model.fit(X_svr_scaled, y_svr)

# Dense grid for drawing a smooth prediction curve, converted back to minutes
X_grid = np.linspace(X_svr_scaled.min(), X_svr_scaled.max(), 300).reshape(-1, 1)
y_grid_pred = svr_model.predict(X_grid)
X_grid_orig = scaler_svr.inverse_transform(X_grid).ravel()

# Support vectors = training points that lie on/outside the epsilon tube

```

```

sv_indices = svr_model.support_
X_sv = X_svr[sv_indices]
y_sv = y_svr[sv_indices]

plt.figure(figsize=(11, 6))
# Training samples
plt.scatter(X_svr, y_svr, color='navy', s=35, alpha=0.75, label='Eruptions')

# Regression surface
plt.plot(X_grid_orig, y_grid_pred, color='crimson', lw=2.5, label='SVR Prediction  $f(x)$ ')

# Epsilon-insensitive tube boundaries + shaded tube
plt.plot(X_grid_orig, y_grid_pred + epsilon_val, color='gray', linestyle='--', lw=1.5, label=rf'Upper Tube ($+\backslash$ epsilon={epsilon_val}$)')
plt.plot(X_grid_orig, y_grid_pred - epsilon_val, color='gray', linestyle='--', lw=1.5, label=rf'Lower Tube ($-\backslash$ epsilon={epsilon_val}$)')
plt.fill_between(X_grid_orig, y_grid_pred - epsilon_val, y_grid_pred + epsilon_val, color='gray', alpha=0.15)

# Support vectors are highlighted - only these points influence the fitted curve
plt.scatter(X_sv, y_sv, s=120, facecolors='none', edgecolors='gold', linewidths=2.5, label=f'Support Vectors ({len(X_sv)})')

plt.title(f"Support Vector Regression on Old Faithful (C=10,  $\epsilon$ ={{epsilon_val}})", fontsize=14, fontweight='bold')
plt.xlabel("Waiting time before eruption (minutes)", fontsize=12)
plt.ylabel("Eruption duration (minutes)", fontsize=12)
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()

# ---- Quantitative evaluation ----
print(f"SVR R2 Score: {r2_score(y_svr, svr_model.predict(X_svr_scaled)):.4f}")
print(f"SVR RMSE: {root_mean_squared_error(y_svr, svr_model.predict(X_svr_scaled)):.4f} minutes")
print(f"Support Vectors: {len(X_sv)} out of {len(X_svr)} points ({len(X_sv)/len(X_svr)*100:.1f}%)")

```

Listing 78: Real 1D regression: Old Faithful geyser

## **Part X**

# **Generative / Foundation Models: Large Language Models**

## Experiment 18: Large Language Model (LLM) Experiment

**Dataset:** SMS Spam Collection (`data/sms_spam.csv`, 5,572 messages; balanced subset 400 train / 200 test) + DistilBERT / DistilGPT2

**Key parameters:** Fine-tune: fresh 2-class head, 2 epochs, AdamW learning rate  $5 \times 10^{-5}$ , batch 16, max length 64

**Headline result:** Held-out accuracy 0.9600; spam F1 0.9592; greedy vs. nucleus generation demonstrated

### Objective

Demonstrate a reproducible LLM workflow for text classification and generation while understanding tokenization, inference and evaluation.

### Suggested Dataset

DistilBERT (SST-2) for tokenization and sentiment, DistilGPT2 for generation, and the real SMS Spam Collection (`data/sms_spam.csv`) for fine-tuning. A balanced subset (300 ham + 300 spam) is split into 400 training and 200 stratified test messages.

### Required Libraries

transformers, torch, pandas.

### Brief Theory

Language models operate on **subword tokens** (Byte-Pair Encoding / WordPiece) mapped to integer vocabulary IDs; an **attention mask** marks real tokens versus padding. Self-attention computes

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}} + \text{mask}\right) V,$$

so every token attends to every other token. Causal generation predicts one token at a time from  $P(t_n | t_1, \dots, t_{n-1}) = \text{softmax}(z_n/T)$ : **greedy** decoding takes the argmax each step, **temperature**  $T$  scales the logits, and **top- $p$**  (nucleus) sampling restricts sampling to the smallest token set whose cumulative probability exceeds  $p$ . **Fine-tuning** reuses pretrained weights, replaces the classification head and trains briefly with a small learning rate.

### Procedure

1. Tokenize a sentence and inspect the tokens, IDs and attention mask.
2. Run the sentiment pipeline on fixed test sentences.
3. Generate text with greedy decoding and with temperature/top- $p$  sampling.
4. Load the SMS Spam Collection and build a balanced subset (300 ham + 300 spam) with a 400/200 stratified split.
5. Fine-tune DistilBERT (fresh 2-class head, AdamW learning rate  $5 \times 10^{-5}$ , 2 epochs, batch 16).
6. Evaluate accuracy and F1 on the held-out messages and inspect example predictions.

### Reference Implementation

```
classifier = pipeline("sentiment-analysis", model=model_id,
                      device=0 if torch.cuda.is_available() else -1)

ft_model = AutoModelForSequenceClassification.from_pretrained(model_id, num_labels=2,
                                                             ignore_mismatched_sizes=True)

optimizer = torch.optim.AdamW(ft_model.parameters(), lr=5e-5)
outputs = ft_model(input_ids=ids, attention_mask=mask, labels=labels)
outputs.loss.backward(); optimizer.step()
```

**Listing 79:** Inference pipeline and the fine-tuning loop

## Evaluation / Expected Output

**Table 23:** DistilBERT fine-tuned on the SMS Spam Collection (200 held-out messages).

| Item                                     | Value                     |
|------------------------------------------|---------------------------|
| Fine-tune loss (epoch 1 $\rightarrow$ 2) | 0.705 $\rightarrow$ 0.079 |
| Held-out accuracy                        | <b>0.9600</b>             |
| Held-out F1 (spam)                       | <b>0.9592</b>             |

Example predictions (all correct): a prize-notification spam at 99.4% confidence, a short ham message at 99.5%, and another spam at 89.7%. Generation: greedy decoding produced the safe deterministic continuation “*the way we think about the world*”, while nucleus sampling ( $T = 0.7$ ,  $p = 0.9$ ) produced a different but coherent sentence — the effect of temperature and top- $p$ .

The experiment demonstrates the complete practical pipeline and shows that fine-tuning a small pretrained model on only 400 real messages already gives 96% spam detection. Limitations: the balanced subset ignores the natural 13% spam prior of the full collection, the models are small and can hallucinate, and prompts/data must be fixed for reproducibility.

## Viva / Discussion Questions

**What is tokenization?** Splitting text into subword units that map to integer vocabulary IDs, letting a fixed vocabulary represent any word; special tokens and attention masks handle sequence framing and padding.

**Pretraining vs fine-tuning vs prompting?** Pretraining learns general language representation from massive corpora; fine-tuning adapts those weights to a task with labeled data; prompting steers a frozen model through instructions or examples.

**Why is LLM output not automatically ground truth?** The model generates the most plausible continuation, not verified facts; it can hallucinate, and confidence scores are not calibrated correctness probabilities.

**How can prompt wording affect results?** Meaning, tone, format and label wording all change the conditional distribution the model samples from, so prompts must be fixed and reported.

## Complete Code Listing

```
# ---- Core numerical and plotting libraries ----
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ---- PyTorch; used for the fine-tuning demo ----
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
# ---- Dataset splitting + metrics for the fine-tune evaluation ----
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
# ---- Hugging Face Transformers: tokenizers, models, pipelines, generation configs ----
from transformers import (
    AutoTokenizer,
    AutoModelForSequenceClassification,
    AutoModelForCausalLM,
    GenerationConfig,
    pipeline
)

# Styling setup (consistent look across all notebooks)
plt.style.use('seaborn-v0_8-whitegrid' if 'seaborn-v0_8-whitegrid' in plt.style.available else 'default')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['figure.figsize'] = (10, 5)
plt.rcParams['figure.dpi'] = 110

torch.manual_seed(42)
torch.cuda.manual_seed_all(42) # reproducible GPU training
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu') # GPU if available
```

```
print(f"PyTorch on {device}")
```

**Listing 80:** Core numerical and plotting libraries

```
# ---- Load a pre-trained DistilBERT tokenizer (fast, subword/WordPiece based) ----
model_id = "distilbert/distilbert-base-uncased-finetuned-sst-2-english"
print(f>Loading Tokenizer: {model_id}...")
tokenizer = AutoTokenizer.from_pretrained(model_id)

sample_text = "Machine learning and deep neural networks are transforming modern science!"
# Tokenize with padding/truncation so the output has a fixed shape (here 20 tokens)
encoded = tokenizer(
    sample_text,
    padding="max_length",
    max_length=20,
    truncation=True,
    return_tensors="pt"
)

# Convert integer IDs back to readable subword tokens
tokens = tokenizer.convert_ids_to_tokens(encoded['input_ids'][0])
print(f>Original Text: {sample_text}\n")
print(f>Subword Tokens ({len(tokens)}):\n", tokens)
print(f>Input IDs: \n", encoded['input_ids'][0].tolist())
print(f>Attention Mask: \n", encoded['attention_mask'][0].tolist())
```

**Listing 81:** Load a pre-trained DistilBERT tokenizer (fast, subword/WordPiece based)

```
# ---- Load a pre-trained sentiment-analysis pipeline (DistilBERT SST-2) ----
classifier = pipeline("sentiment-analysis", model=model_id, device=0 if torch.cuda.is_available() else -1)

test_sentences = [
    "The new neural network model demonstrated extraordinary accuracy and remarkable speed.",
    "The dataset was poorly labeled, noisy, and the model completely failed to converge.",
    "The gradient descent algorithm performed reasonably well, meeting standard expectations.",
    "A catastrophic failure occurred during training due to exploding gradient issues."
]

# One forward pass per sentence -> label + calibrated confidence score
predictions = classifier(test_sentences)

# Collect results in a table for display
results_table = []
for sent, pred in zip(test_sentences, predictions):
    results_table.append({
        "Sentence": sent,
        "Predicted Label": pred['label'],
        "Confidence Score": f"{pred['score']*100:.2f}%"
    })

df_preds = pd.DataFrame(results_table)
display(df_preds)
```

**Listing 82:** Load a pre-trained sentiment-analysis pipeline (DistilBERT SST-2)

```
# ---- Load a small causal LM (DistilGPT2) for text generation ----
gen_model_id = "distilbert/distilgpt2"
print(f>Loading Text Generator: {gen_model_id}...")
# clean_up_tokenization_spaces=False keeps BPE spacing intact
generator = pipeline("text-generation", model=gen_model_id,
    device=0 if torch.cuda.is_available() else -1, clean_up_tokenization_spaces=False)

prompt = "Artificial Intelligence will fundamentally transform"

# 1. Greedy Search: always pick the most likely next token (deterministic)
greedy_output = generator(
    prompt, generation_config=GenerationConfig(max_new_tokens=40, do_sample=False)
)[0]['generated_text']

# 2. Temperature + Top-p (Nucleus) Sampling: sample from the most likely tokens up to cumulative p
sampled_output = generator(
    prompt, generation_config=GenerationConfig(max_new_tokens=40, do_sample=True, temperature=0.7, top_p=0.9)
)[0]['generated_text']
```



```
print(f"PROMPT: '{prompt}'\n")
print(f"--- 1. GREEDY SEARCH OUTPUT ---\n{greedy_output.strip()}\n")
print(f"--- 2. NUCLEUS SAMPLING (T=0.7, p=0.9) OUTPUT ---\n{sampling_output.strip()}\n")
```

Listing 83: Load a small causal LM (DistilGPT2) for text generation

```
# ---- Real dataset: SMS Spam Collection (5,572 labeled messages) ----
sms = pd.read_csv('../data/sms_spam.csv')
sms['y'] = (sms['label'] == 'spam').astype(int)
print(f"SMS Spam Collection: {len(sms)} messages | ham {(sms.y == 0).sum()}, spam {(sms.y == 1).sum()}")

# Balanced subset for a fast fine-tune demo: 300 ham + 300 spam
ham = sms[sms.y == 0].sample(n=300, random_state=42)
spam = sms[sms.y == 1].sample(n=300, random_state=42)
subset = pd.concat([ham, spam]).sample(frac=1, random_state=42).reset_index(drop=True)

# Stratified split keeps the ham/spam ratio identical in train and test
train_df, test_df = train_test_split(subset, test_size=200, random_state=42, stratify=subset.y)
print(f"Train messages: {len(train_df)} | Test messages: {len(test_df)} | spam share: {subset.y.mean():.1%}")

class SMSDataset(Dataset):
    """Tokenize SMS texts and expose them as PyTorch samples."""
    def __init__(self, texts, labels, tokenizer):
        self.encodings = tokenizer(list(texts), padding=True, truncation=True, max_length=64, return_tensors="pt")
        self.labels = list(labels)

    def __getitem__(self, idx):
        item = {key: val[idx] for key, val in self.encodings.items()}
        item['labels'] = torch.tensor(self.labels[idx], dtype=torch.long)
        return item

    def __len__(self):
        return len(self.labels)

train_loader = DataLoader(SMSDataset(train_df.text, train_df.y, tokenizer), batch_size=16, shuffle=True)
test_loader = DataLoader(SMSDataset(test_df.text, test_df.y, tokenizer), batch_size=32)

# ---- Load the pre-trained classifier with a fresh 2-class head ----
# ignore_mismatched_sizes allows replacing the original SST-2 head with a new one.
ft_model = AutoModelForSequenceClassification.from_pretrained(model_id, num_labels=2, ignore_mismatched_sizes=True)
ft_model.to(device)
optimizer = torch.optim.AdamW(ft_model.parameters(), lr=5e-5) # small LR for fine-tuning

print("Pre-trained Transformer loaded for fine-tuning!")
```

Listing 84: Real dataset: SMS Spam Collection (5,572 labeled messages)

```
# ---- Fine-tuning loop (2 epochs on 400 real SMS messages) ----
epochs = 2
ft_model.train()
epoch_losses = []

for epoch in range(epochs):
    running_loss = 0.0
    for batch in train_loader:
        optimizer.zero_grad()

        # Move the batch to the selected device
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['labels'].to(device)

        # Forward pass: the model computes the classification loss internally
        outputs = ft_model(input_ids=input_ids, attention_mask=attention_mask, labels=labels)
        loss = outputs.loss
        loss.backward() # backpropagate through the whole transformer
        optimizer.step() # update weights

    running_loss += loss.item() * len(input_ids) # de-average the batch loss

avg_loss = running_loss / len(train_df)
epoch_losses.append(avg_loss)
print(f"Epoch [{epoch+1}/{epochs}] Fine-Tuning Loss: {avg_loss:.4f}")
```

**Listing 85:** Fine-tuning loop (2 epochs on 400 real SMS messages)

```

# ---- Evaluate the fine-tuned model on held-out real SMS messages ----
ft_model.eval()
preds, targets = [], []

with torch.no_grad():
    for batch in test_loader:
        # Move the batch to the same device as the fine-tuned model (GPU when available)
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        logits = ft_model(input_ids=input_ids, attention_mask=attention_mask).logits
        preds += logits.argmax(-1).tolist()      # predicted class (0 = ham, 1 = spam)
        targets += batch['labels'].tolist()

print(f"Held-out test accuracy: {accuracy_score(targets, preds):.4f}")
print(f"Held-out test F1 (spam): {f1_score(targets, preds):.4f}")

# ---- Show a few individual predictions with confidence ----
labels_map = {0: "ham", 1: "spam"}
with torch.no_grad():
    for i in range(3):
        message = test_df.text.iloc[i]
        enc = tokenizer(message, truncation=True, max_length=64, return_tensors="pt").to(device)
        probs = torch.softmax(ft_model(**enc).logits, dim=1).cpu().numpy()[0]
        print(f"\nMessage: {message[:75]}...")
        print(f"True: {labels_map[test_df.y.iloc[i]]} | Predicted: {labels_map[probs.argmax()]} ({probs.max()
        ↪ *100:.1f}%)")

```

**Listing 86:** Evaluate the fine-tuned model on held-out real SMS messages

## **Part XI**

# **Non-Parametric Regression: Generalized Regression Neural Network**

## Experiment 19: Generalized Regression Neural Network (GRNN)

**Dataset:** Real motorcycle impact data (`data/motorcycle.csv`, `R MASS::mcycle`); 133 observations, time vs head acceleration

**Key parameters:** GRNN from scratch (NumPy); smoothing spread  $\sigma$  selected by 5-fold cross-validation over 0.02–1.0

**Headline result:**  $\sigma = 0.08$  (CV RMSE 25.00 g); held-out  $R^2$  0.7255, RMSE 22.78 g

### Objective

Implement GRNN-style kernel regression from scratch and study the effect of the smoothing parameter  $\sigma$  with cross-validation.

### Suggested Dataset

Real motorcycle impact data (`data/motorcycle.csv`, `R MASS::mcycle`): 133 observations of time after impact (ms) versus head acceleration (g). Time is standardized before computing kernel distances.

### Required Libraries

numpy, pandas, matplotlib, scikit-learn (base classes, KFold, metrics).

*Content, result plots and complete code listings for this experiment are filled in after the format review.*

### Brief Theory

### Procedure

### Reference Implementation

### Evaluation / Expected Output

### Viva / Discussion Questions

### Complete Code Listing

## **Part XII**

### **Integrated Mini Project and Appendices**

## Integrated Mini Project: Comparative Machine Learning Study

**Scope:** Compare algorithms from different families on identical data, splits and seeds

**Protocol:** Fixed random seeds; untouched test sets; two or more metrics per task

**Headline result:** Gradient boosting leads tabular regression; ensembles tie on small classification data; density clustering finds real seismic zones without labels

### Objective

For the final laboratory assessment, compare at least three algorithms from different families on one real-world dataset, with a controlled, explainable experiment rather than a search for the largest score.

### Suggested Dataset

The manual's own experiments provide the comparison material: Heart Disease (classification), California Housing (regression), Mall Customers (clustering) and the USGS earthquake catalogue (density-based clustering).

### Required Libraries

scikit-learn, xgboost, catboost, numpy, pandas, matplotlib.

*Content, result plots and complete code listings for this experiment are filled in after the format review.*

### Workflow

### Comparative Results

### Discussion and Conclusions

### Viva / Discussion Questions

## Appendix A: Quick Reference Formulas

| Concept            | Formula / Definition                                                                  |
|--------------------|---------------------------------------------------------------------------------------|
| Euclidean distance | $d(x, y) = \sqrt{\sum_j (x_j - y_j)^2}$                                               |
| K-means objective  | $J = \sum_i \ x_i - \mu_{c_i}\ ^2$                                                    |
| Regression MAE     | $\text{MAE} = \frac{1}{n} \sum_i  y_i - \hat{y}_i $                                   |
| Regression RMSE    | $\text{RMSE} = \sqrt{\frac{1}{n} \sum_i (y_i - \hat{y}_i)^2}$                         |
| $R^2$              | $R^2 = 1 - \text{SS}_{\text{res}} / \text{SS}_{\text{tot}}$                           |
| Precision          | $\text{TP} / (\text{TP} + \text{FP})$                                                 |
| Recall             | $\text{TP} / (\text{TP} + \text{FN})$                                                 |
| F1                 | $2 \times \text{Precision} \times \text{Recall} / (\text{Precision} + \text{Recall})$ |
| Sigmoid            | $\sigma(z) = 1 / (1 + e^{-z})$                                                        |
| SVM margin idea    | Maximize margin while penalizing violations via $C$                                   |
| GRNN kernel        | $w_i(x) = \exp(-\ x - x_i\ ^2 / (2\sigma^2))$                                         |

## Appendix B: Reproducibility Checklist

- Python version recorded.
- Package versions recorded.
- Random seed recorded where supported (all experiments use seed 42).
- Dataset source and version recorded (public datasets committed under `data/`).
- Preprocessing steps recorded (scaling fitted on training data only; haversine radians for geospatial data).
- Train/validation/test protocol recorded (stratified, chronological, early-stopping split, 5-fold CV).
- Hyperparameters recorded.
- Evaluation metrics defined before final testing.
- Plots and tables saved with meaningful names.
- Code runs from a clean environment or requirements file.



## Appendix C: Suggested Assessment Rubric

| Component                                     | Suggested Weight |
|-----------------------------------------------|------------------|
| Pre-lab preparation / algorithm understanding | 10%              |
| Correct implementation                        | 25%              |
| Experimental design and preprocessing         | 15%              |
| Evaluation and visualization                  | 15%              |
| Analysis and interpretation                   | 20%              |
| Report quality / reproducibility              | 10%              |
| Viva / discussion                             | 5%               |

# Appendix D: Student Experiment Record

| Field                | Student Entry |
|----------------------|---------------|
| Name                 |               |
| ID                   |               |
| Section              |               |
| Experiment No.       |               |
| Date                 |               |
| Dataset              |               |
| Algorithm / Version  |               |
| Key Hyperparameters  |               |
| Random Seed          |               |
| Main Result          |               |
| Observation          |               |
| Instructor Signature |               |